

# How flexible should my algorithms be?

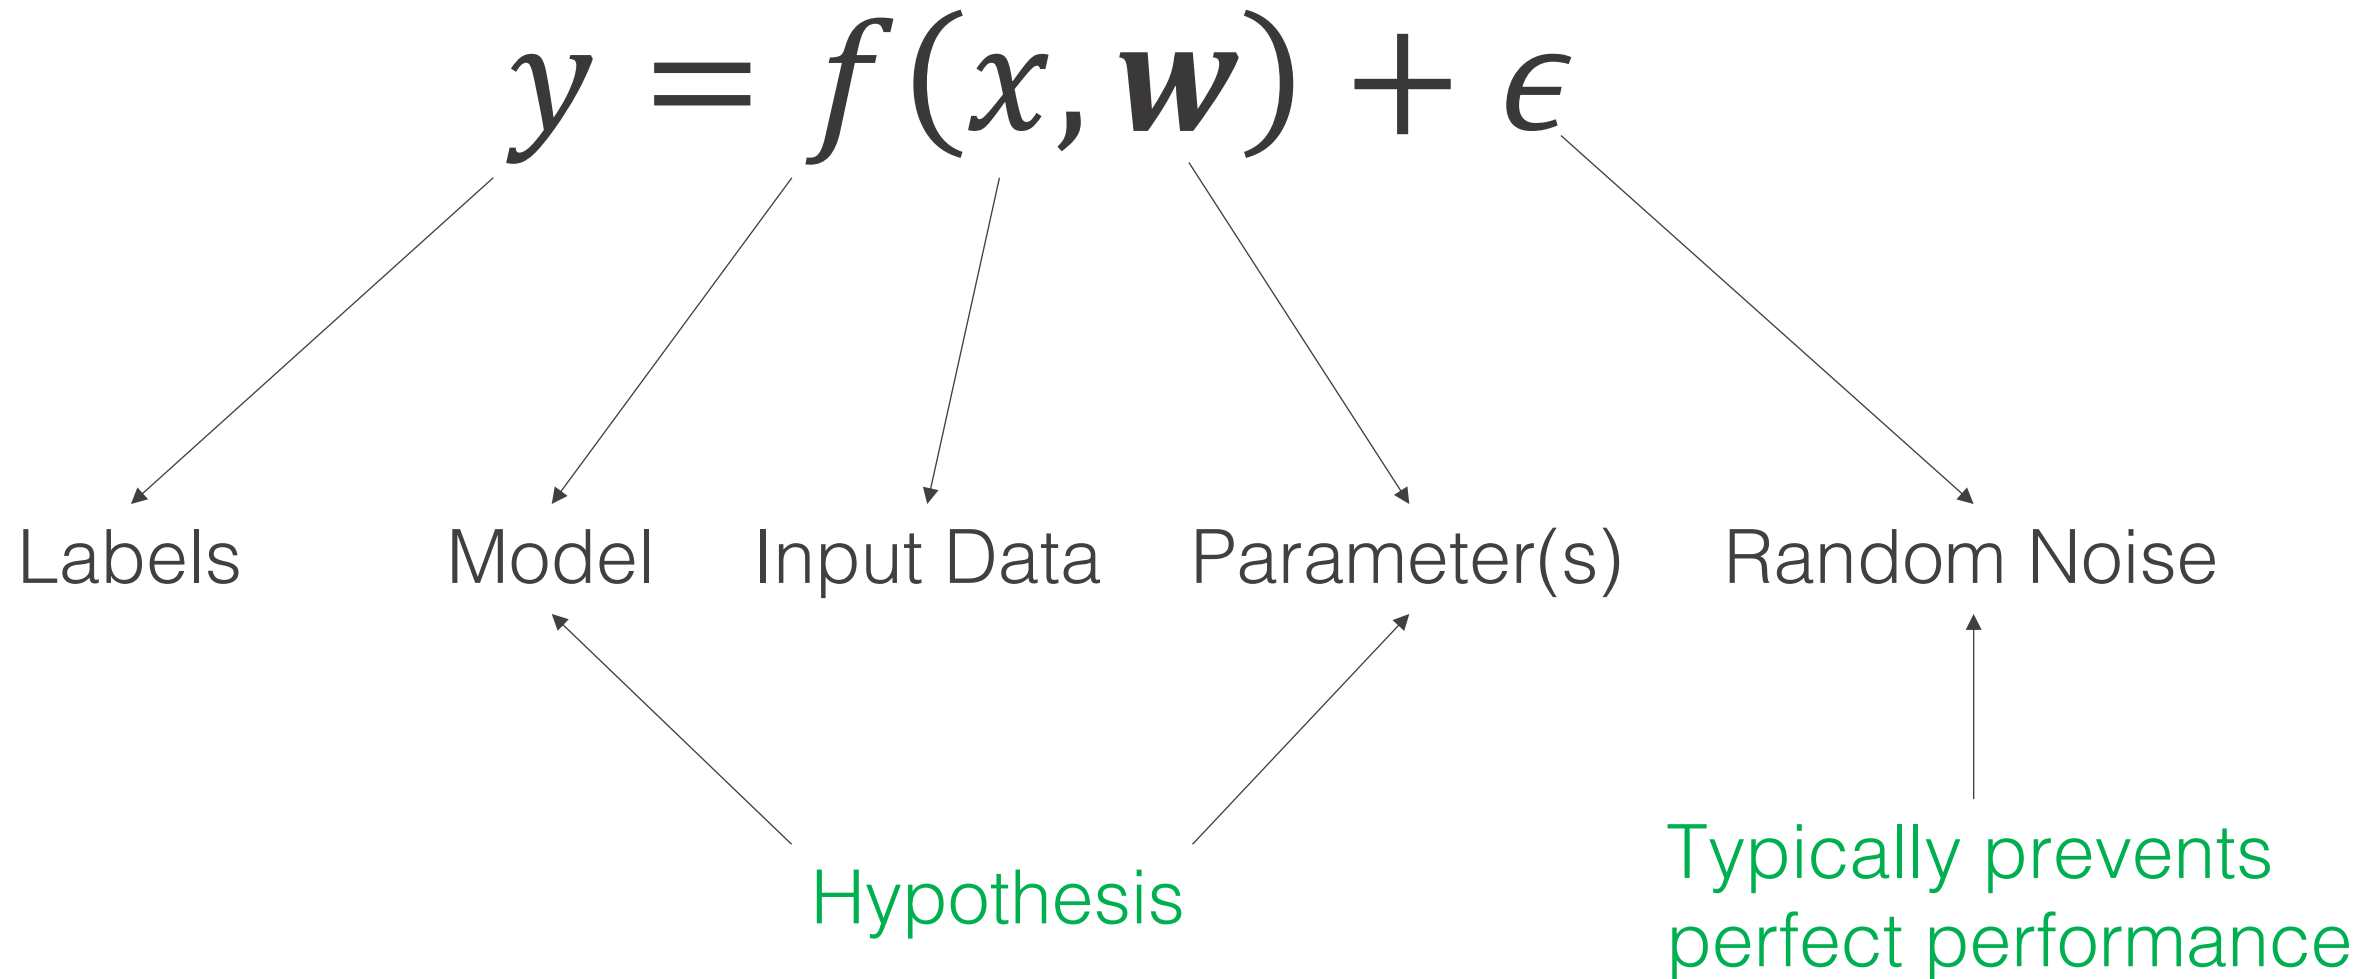
Lecture 03

# Supervised Learning

Algorithm development and application pipeline

# Supervised machine learning model

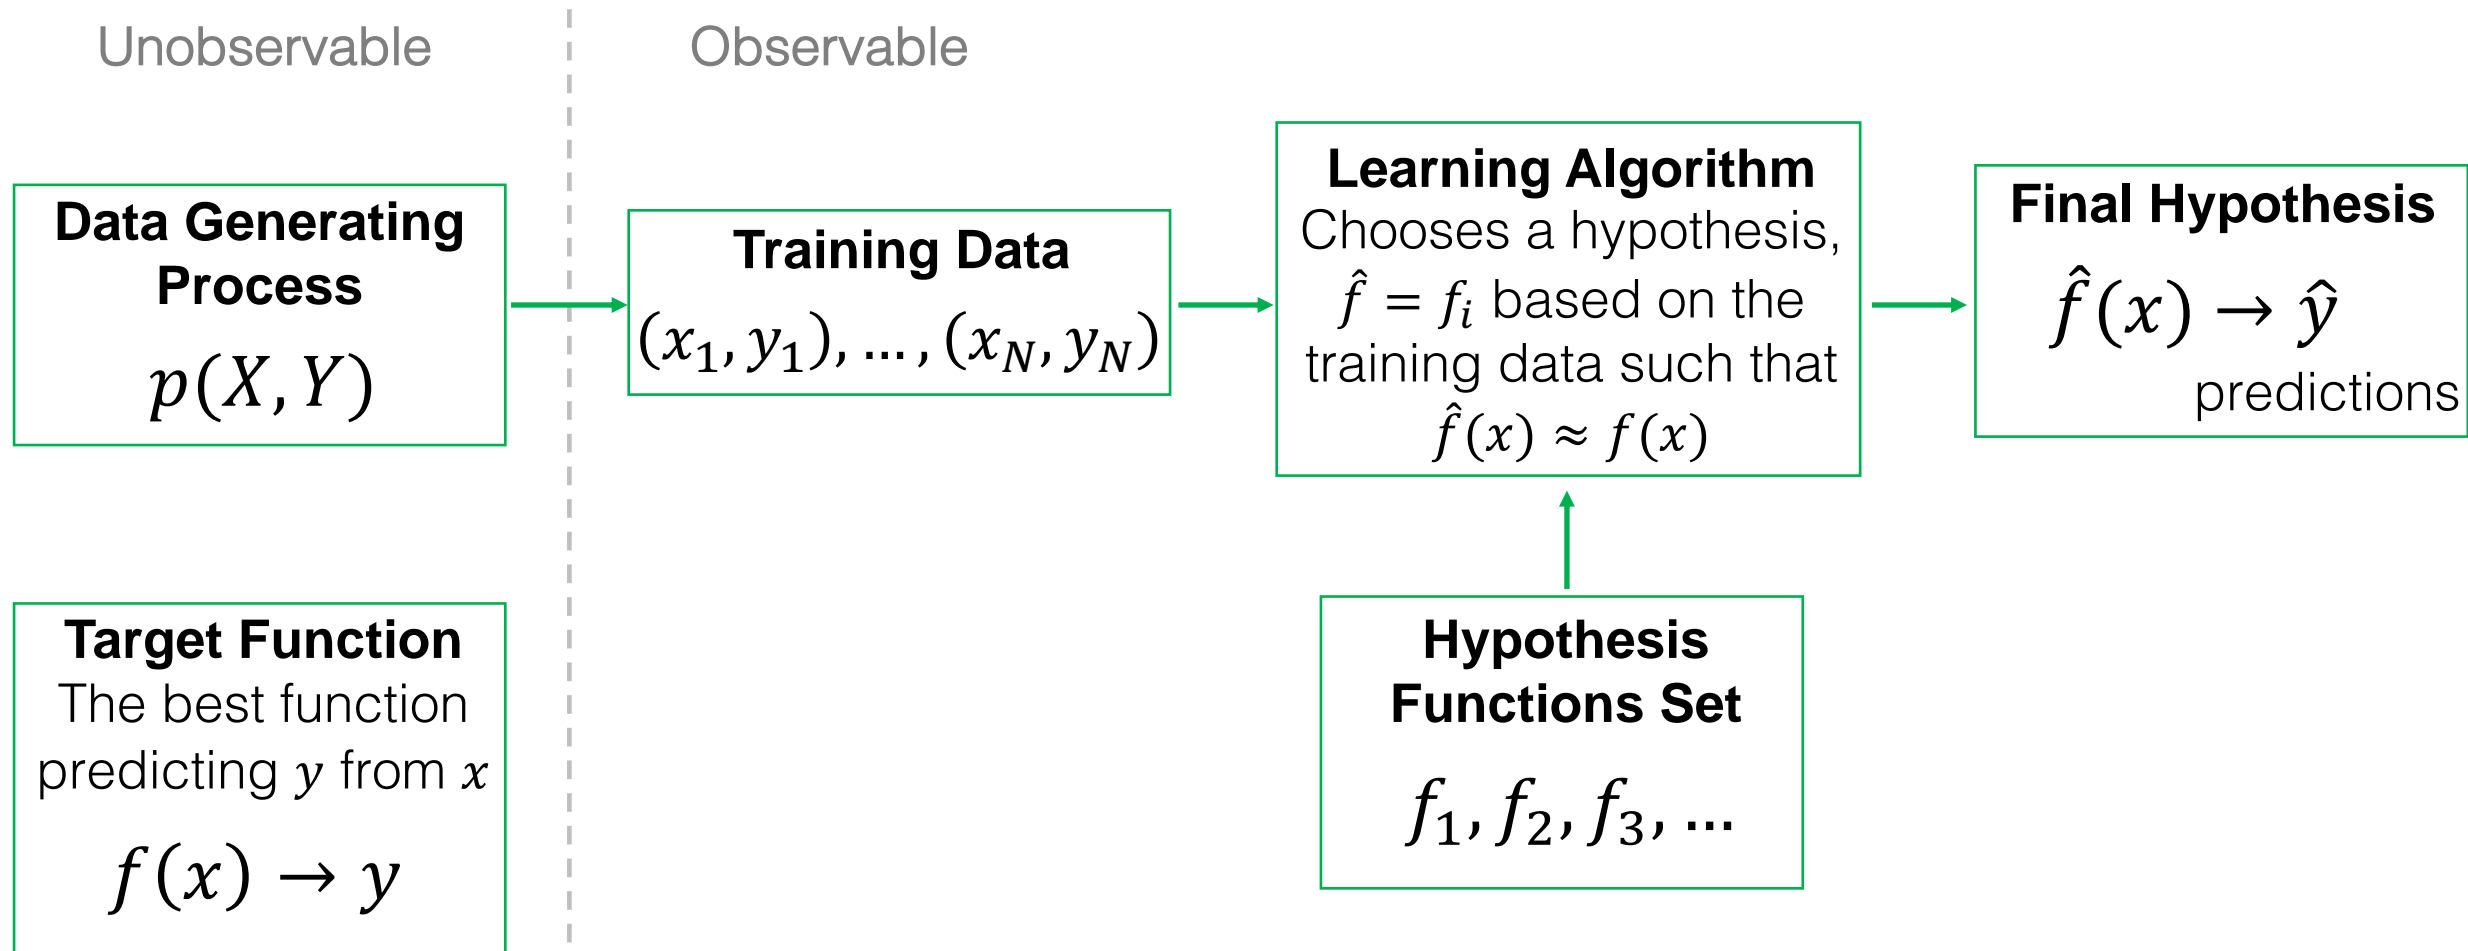
We search for the model that best fits our data



# Components of supervised learning

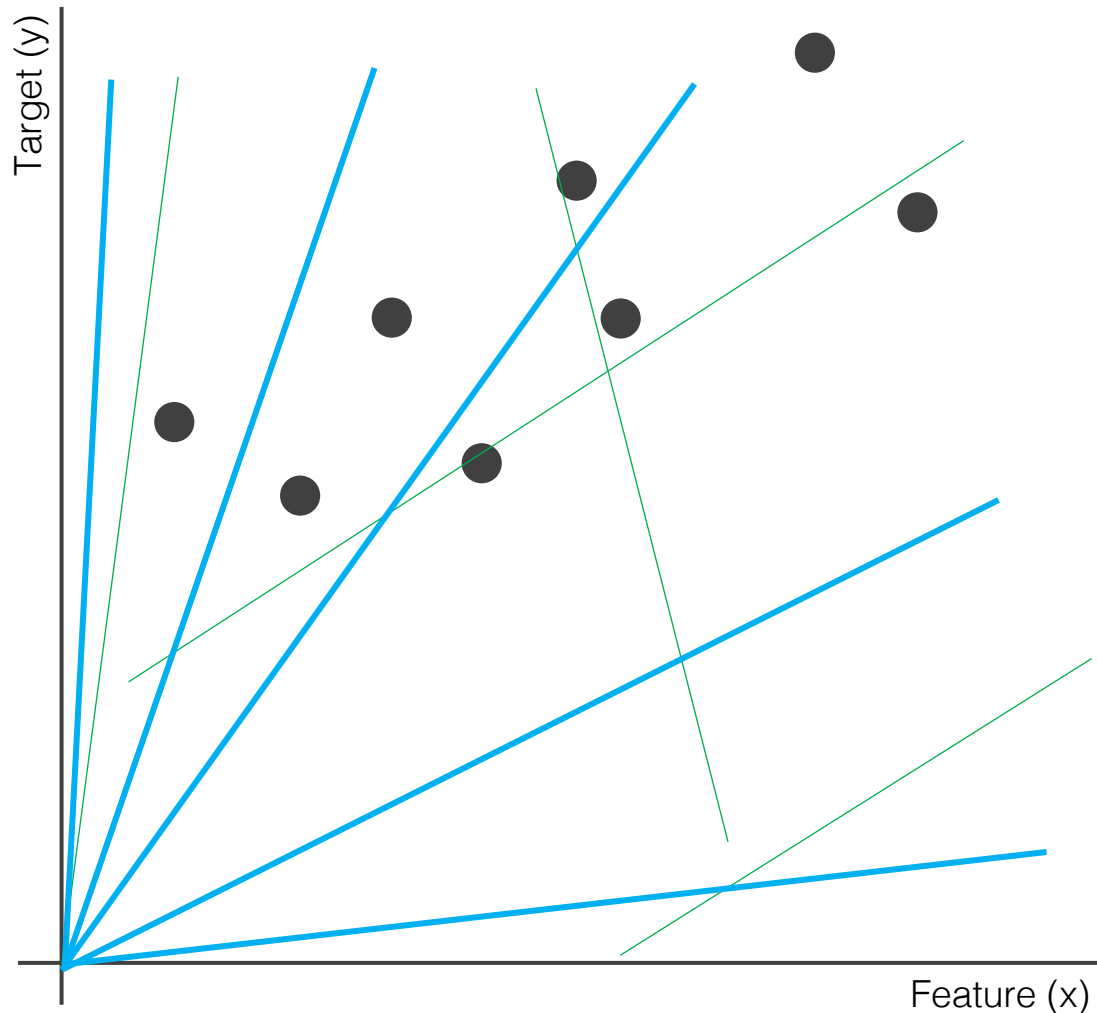
<b>Input</b>	$\mathbf{x}$	
<b>Output</b>	$y$	
<b>Training Data</b>	$(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_N, y_N)$	
<b>Target function</b>	$f(\mathbf{x}) \rightarrow y$	This is unknown, but the best you could ever do
<b>Hypothesis set</b>	$f_i(\mathbf{x}) \rightarrow \hat{y}$	Functions to consider in trying to approximate $f(\mathbf{x})$
<b>Learning algorithm</b>	Optimization technique that searches the hypothesis set for the function $f_i$ that best approximates $f$ (typically by choosing parameters in a model)	

# Supervised Learning



- Need to select the hypothesis functions (models to train)
- Need to select the learning algorithm (for fitting the models to the data)

# Example: linear regression



Using any line as a hypothesis function, how many possible hypothesis functions are in the set?

**Infinitely many**

Using the line  $y = wx$  as the family of hypothesis functions, how many possible hypothesis functions are in the set?

**Infinitely many**

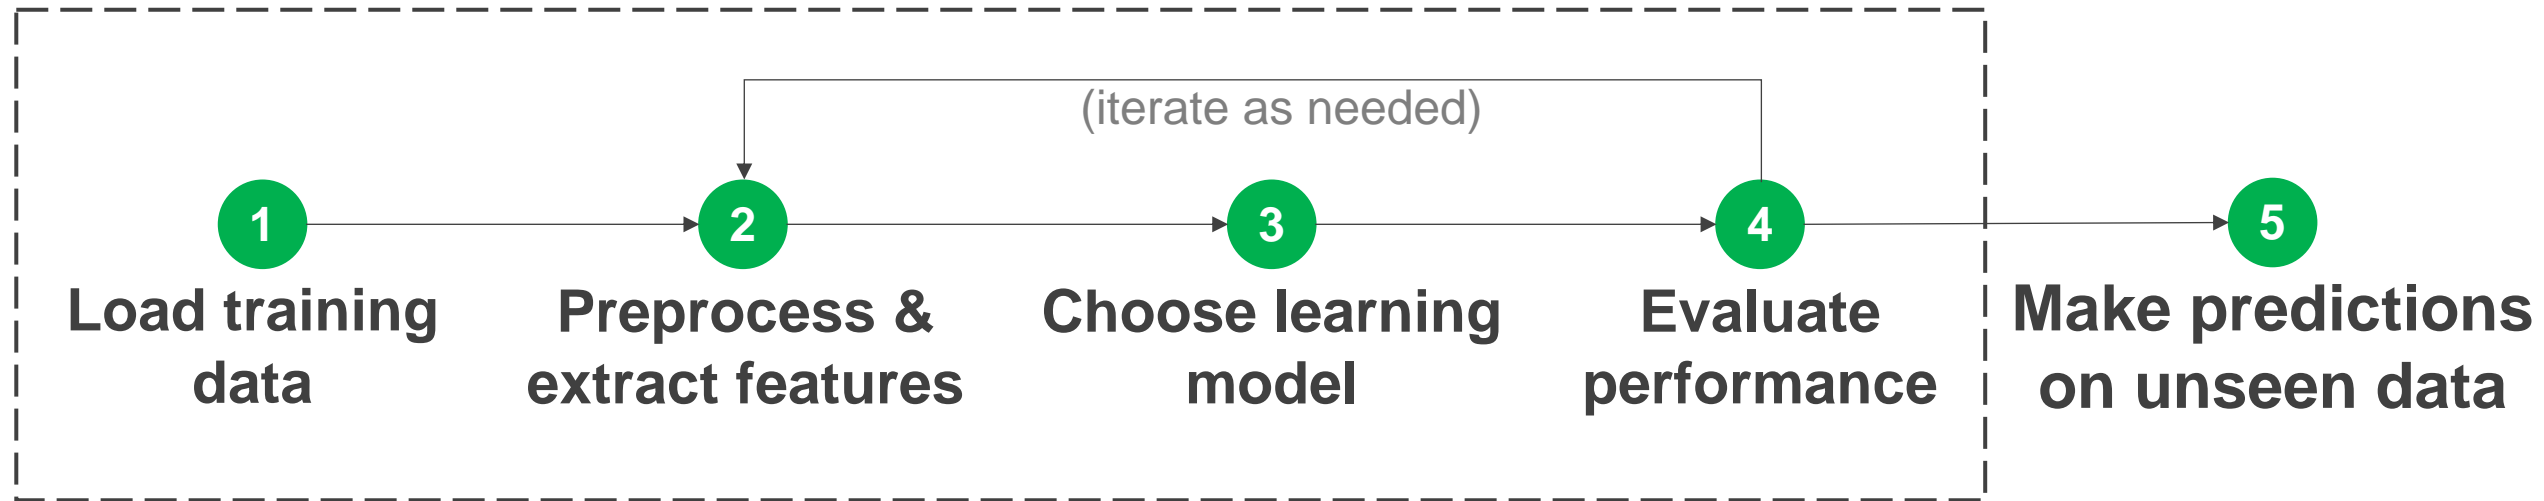
Which set contains the better hypothesis?

Which set has more options to consider?

What is our learning algorithm?

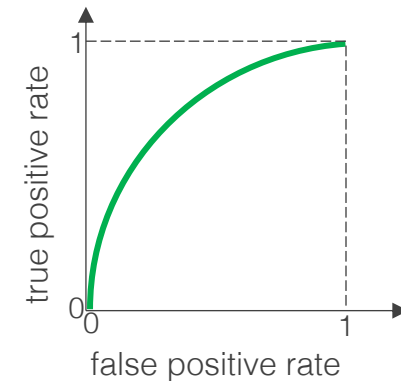
# Algorithm Development

# Application

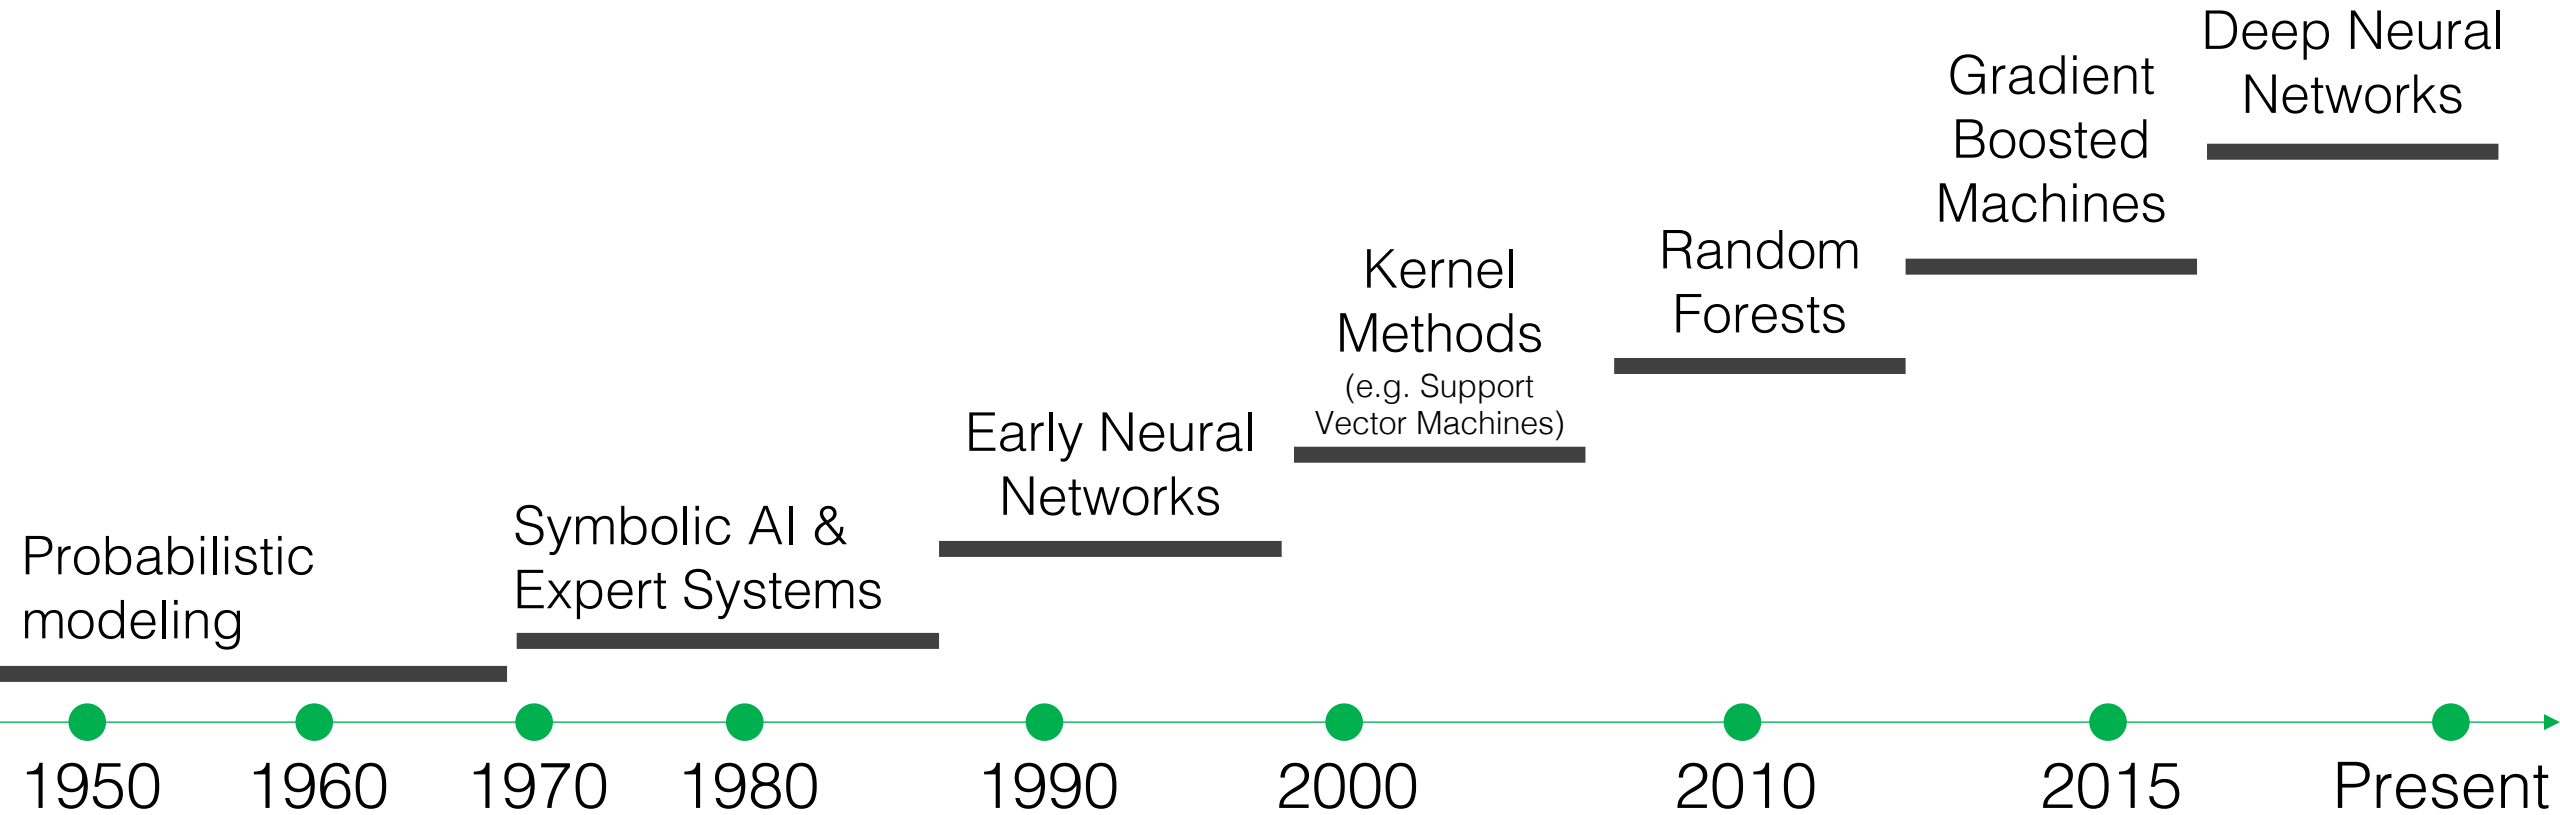


$y$	$X$	$X'$
1	$x_1$	$x_1$ 0.38 0.39 0.85 0.78
1	$x_2$	$x_2$ 0.81 0.91 0.97 0.53
0	$x_3$	$x_3$ 0.65 0.59 0.91 0.11
0	$x_4$	$x_4$ 0.94 0.05 0.40 0.26
1	$x_5$	$x_5$ 0.27 0.19 0.03 0.64
0	$x_6$	$x_6$ 0.02 0.98 0.36 0.11

linear discriminant  
perceptron  
logistic regression  
decision trees  
random forests  
support vector machine  
k nearest neighbors  
neural networks



# Historic Progression of Algorithms



François Chollet, *Deep Learning with Python*, 2017



# How flexible should my algorithms be?

Lecture 03

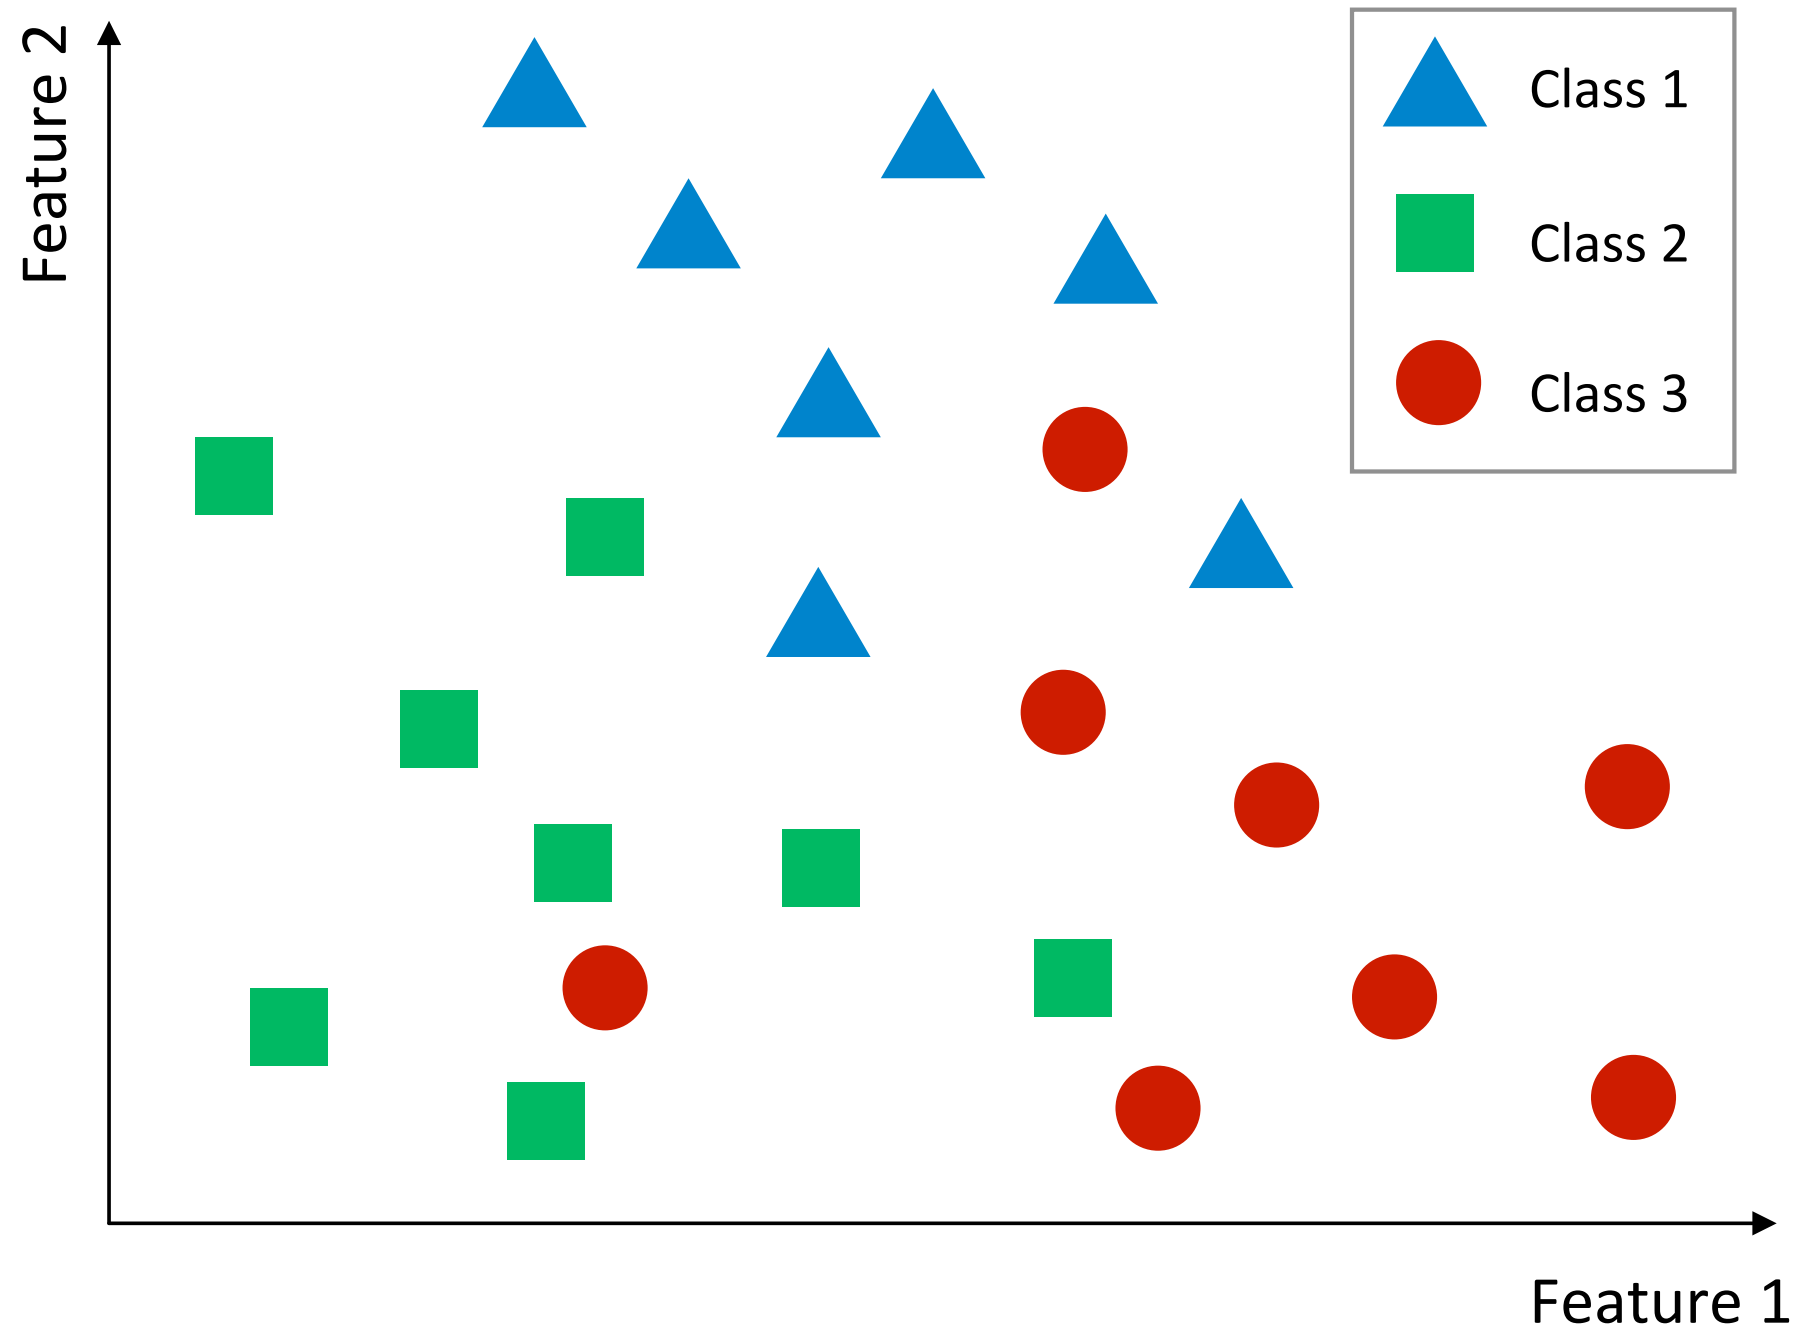
# K-Nearest Neighbors

Classification and Regression

# K Nearest Neighbor Classifier

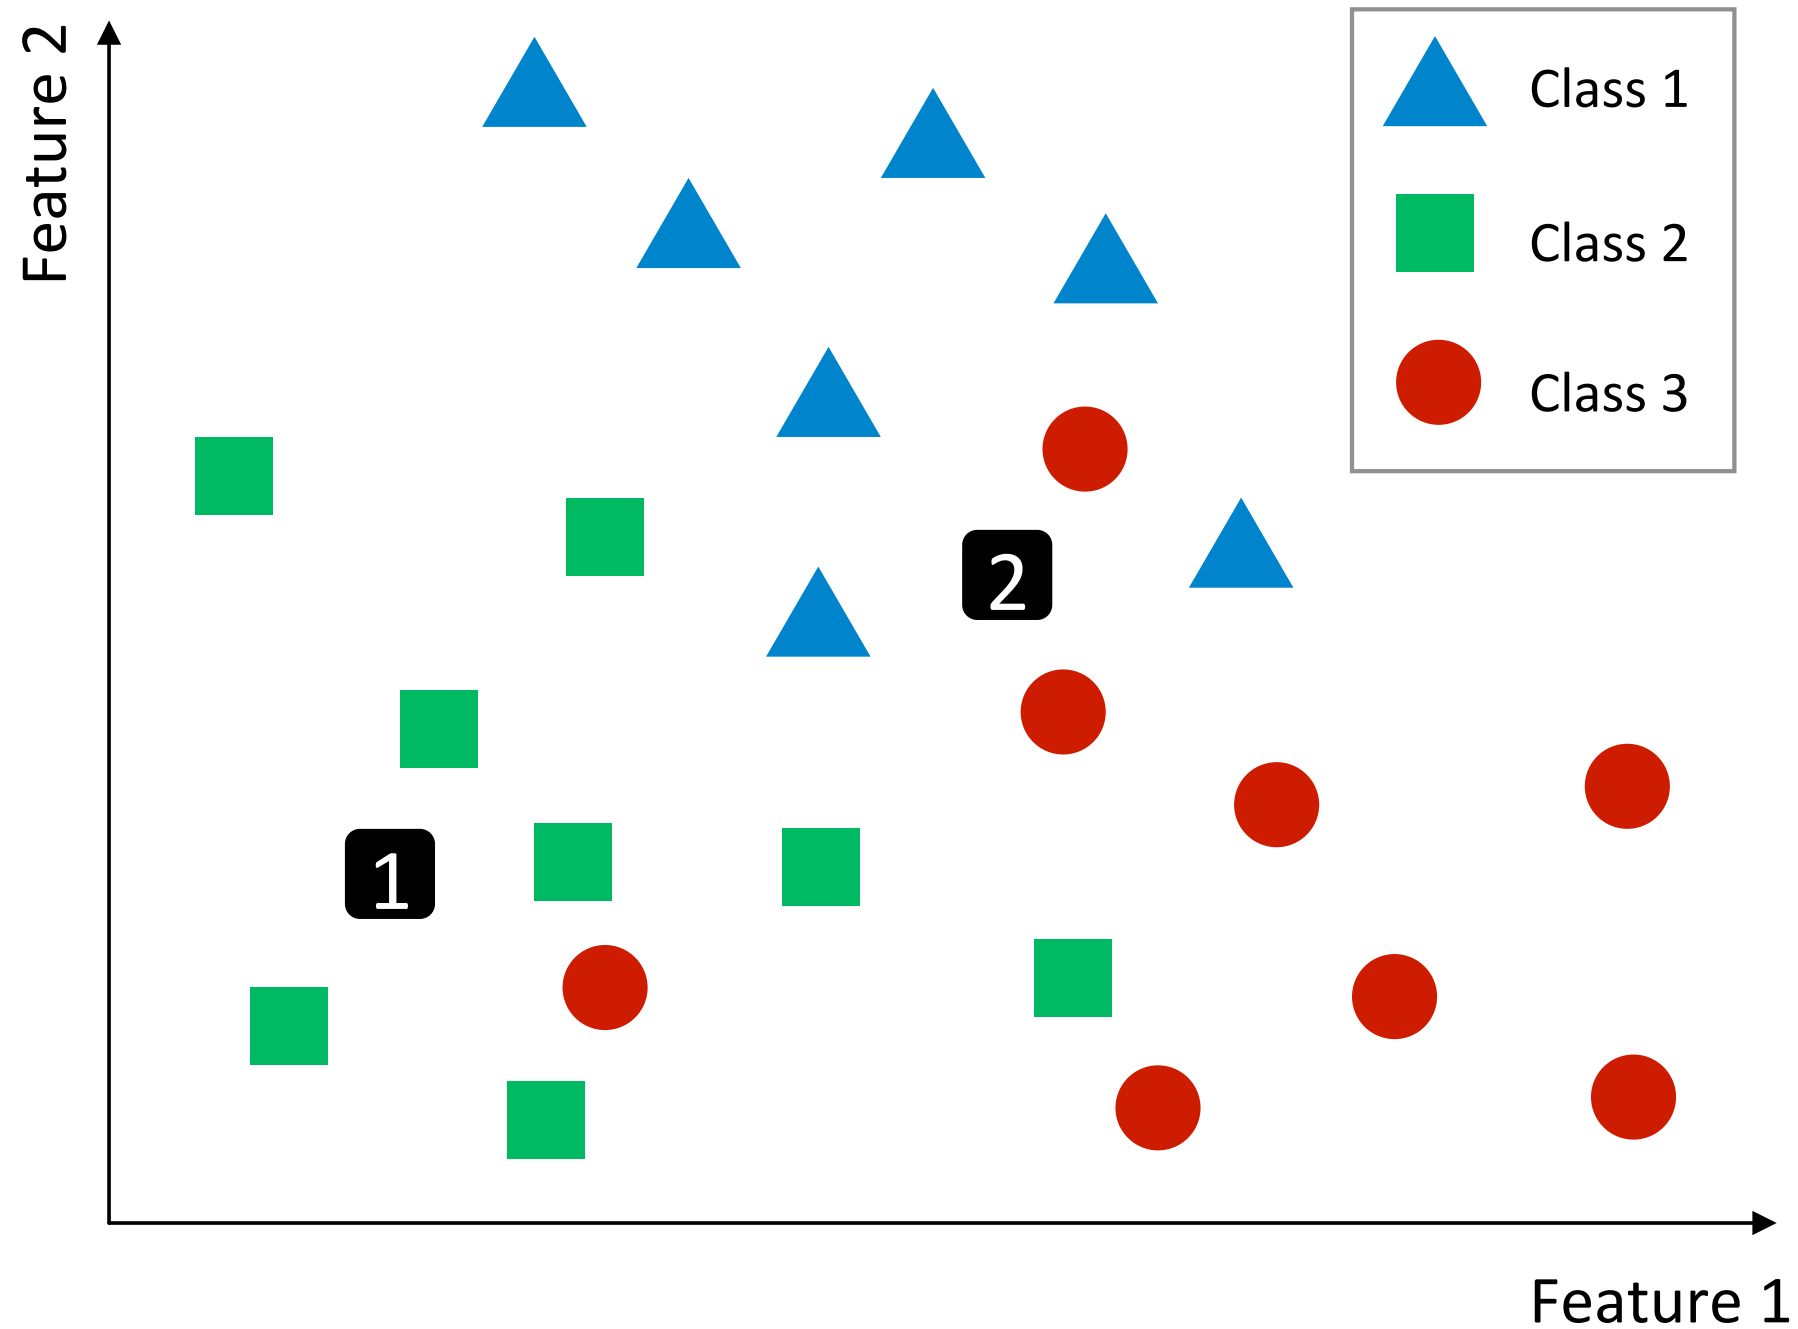
## Step 1: Training

Every new data point is  
a model parameter



# K Nearest Neighbor Classifier

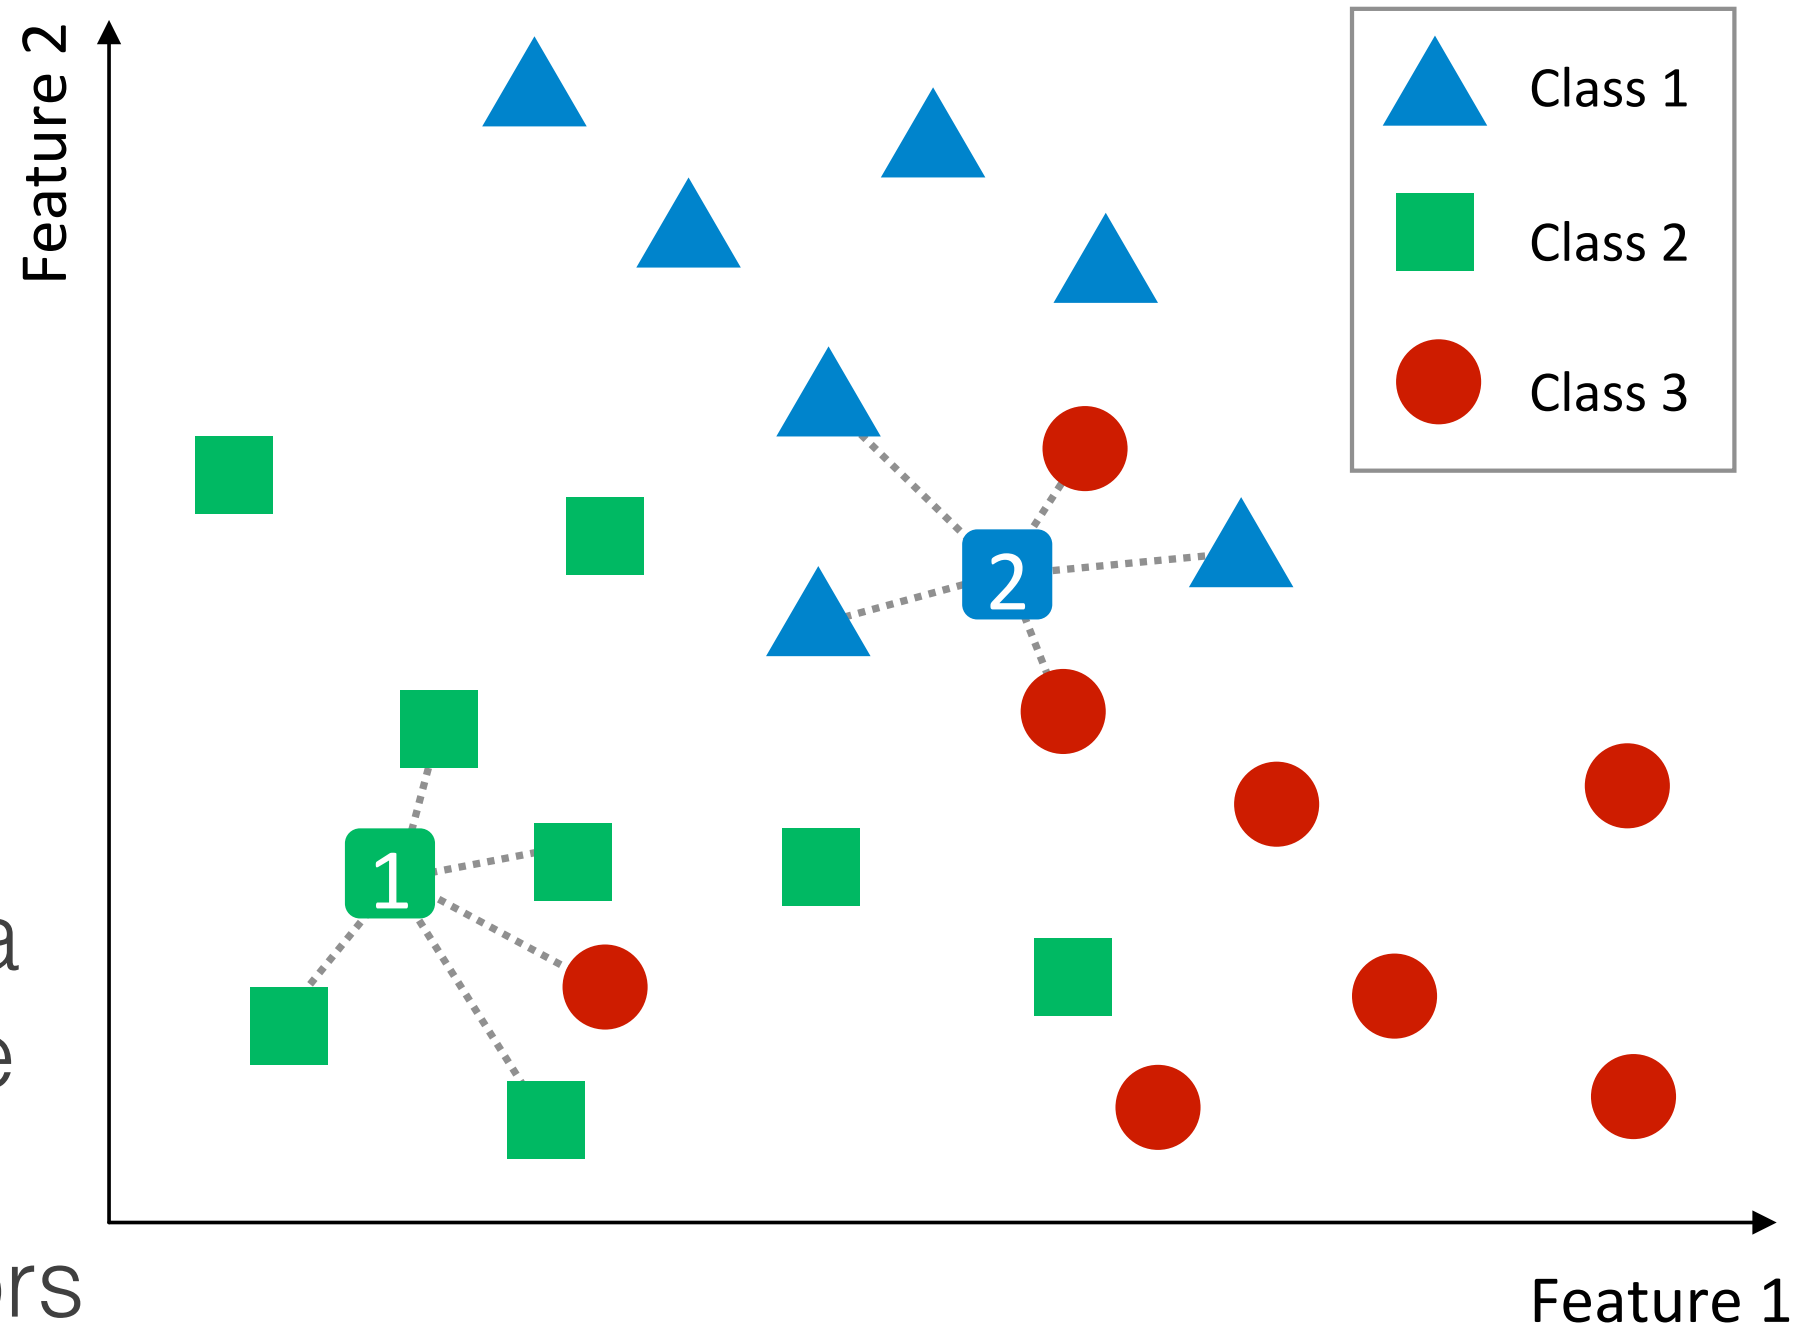
**Step 2:**  
Place new  
(unseen)  
examples in the  
feature space



# K Nearest Neighbor Classifier

## Step 3:

Classify the data by assigning the class of the  $k$  nearest neighbors



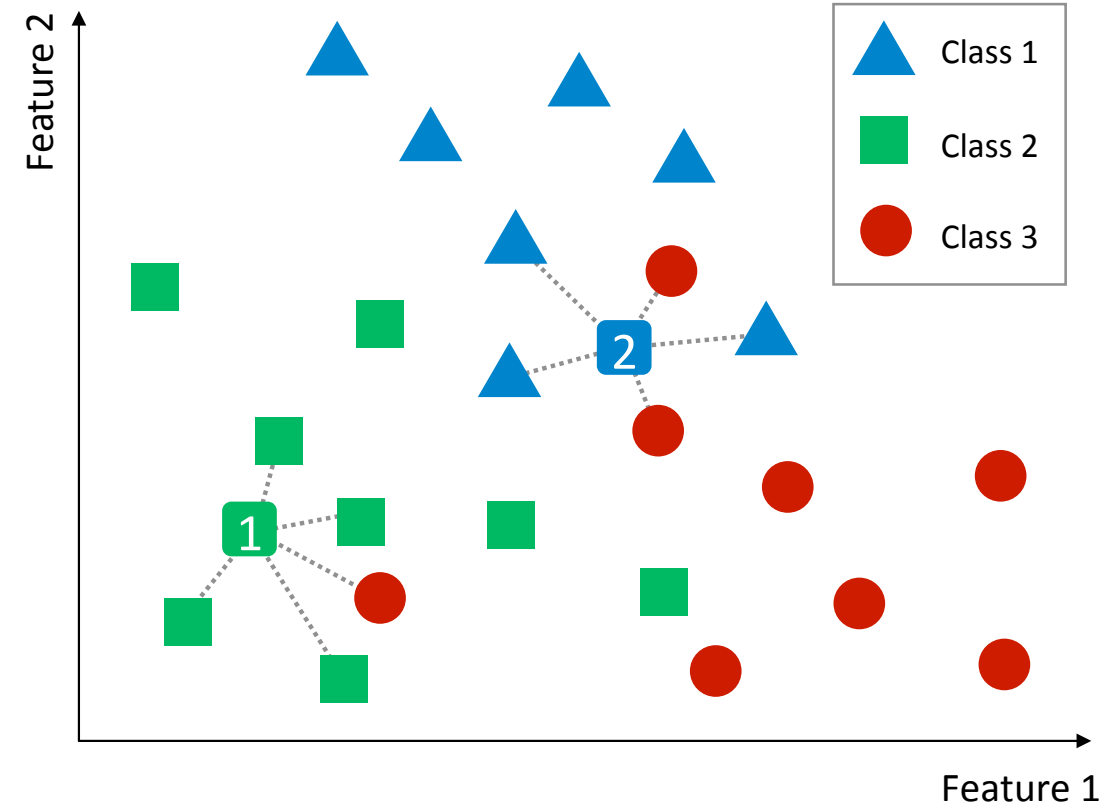
# K Nearest Neighbor Classifier

## Score vs Decision :

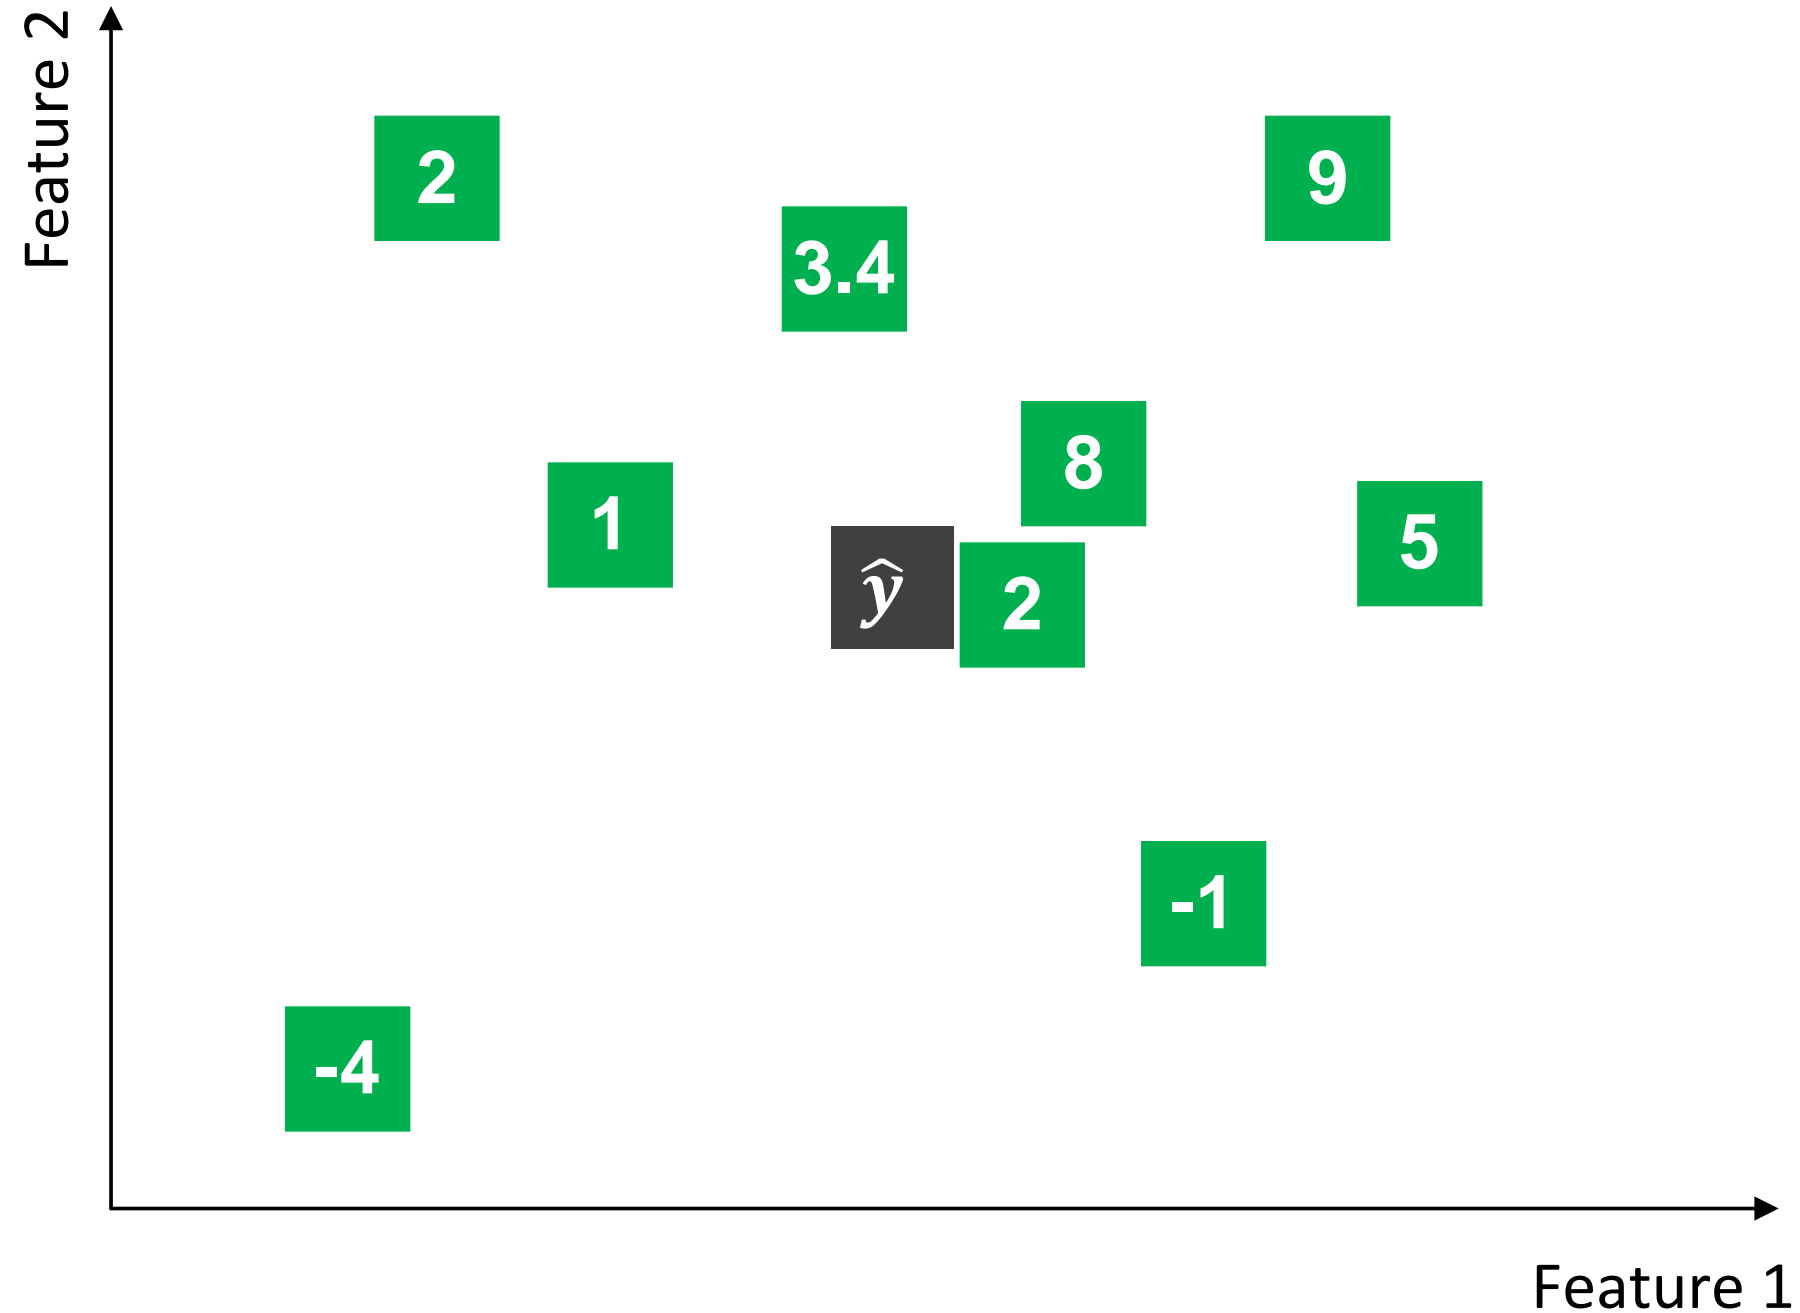
For 5-NN, the confidence score that a sample belongs to a class could be:  $\{0, 1/5, 2/5, 3/5, 4/5, 1\}$

## Decision Rule:

If the confidence score for a class  $>$  threshold, predict that class

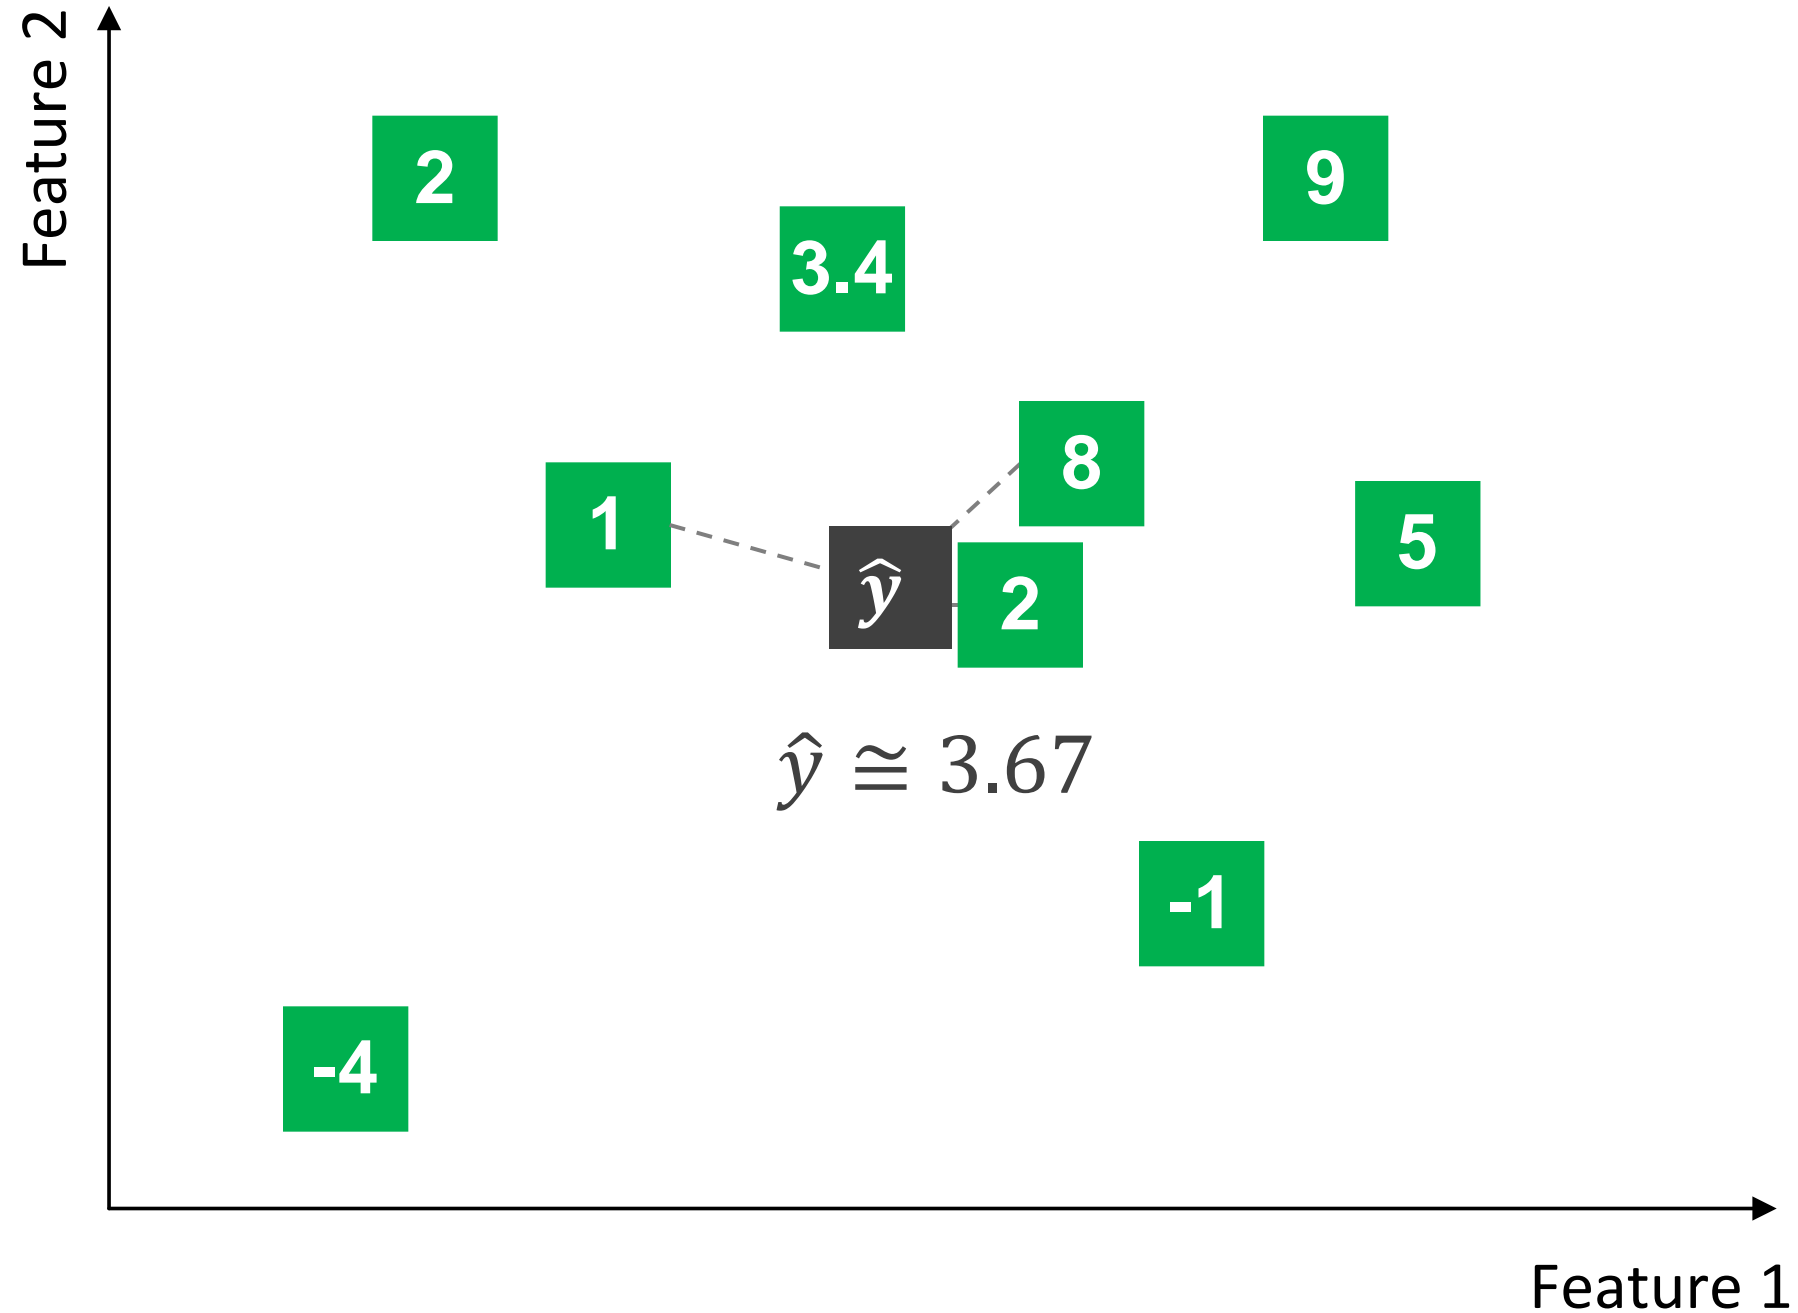


# K Nearest Neighbor Regression



# K Nearest Neighbor Regression

$$\hat{y} = \frac{1}{k} \sum_{y_i \in \{\text{k nearest}\}} y_i$$





# KNN Pros and Cons

## Pros

- Simple to implement and interpret
- Minimal training time
- Naturally handles multiclass data

## Cons

- Computational expensive to find nearest neighbors
- Requires all of the training data to be stored in the model
- Suffers if classes are imbalanced
- Performance may suffer in high dimensions

# How flexible should my model be?

the bias-variance tradeoff and learning to generalize

## **bias**

consistently incorrect  
prediction

error from poor model assumptions  
(high bias results in underfit)

## **variance**

inconsistent prediction

error from sensitivity to  
small changes in the training data  
(high variance results in overfit)

## **noise**

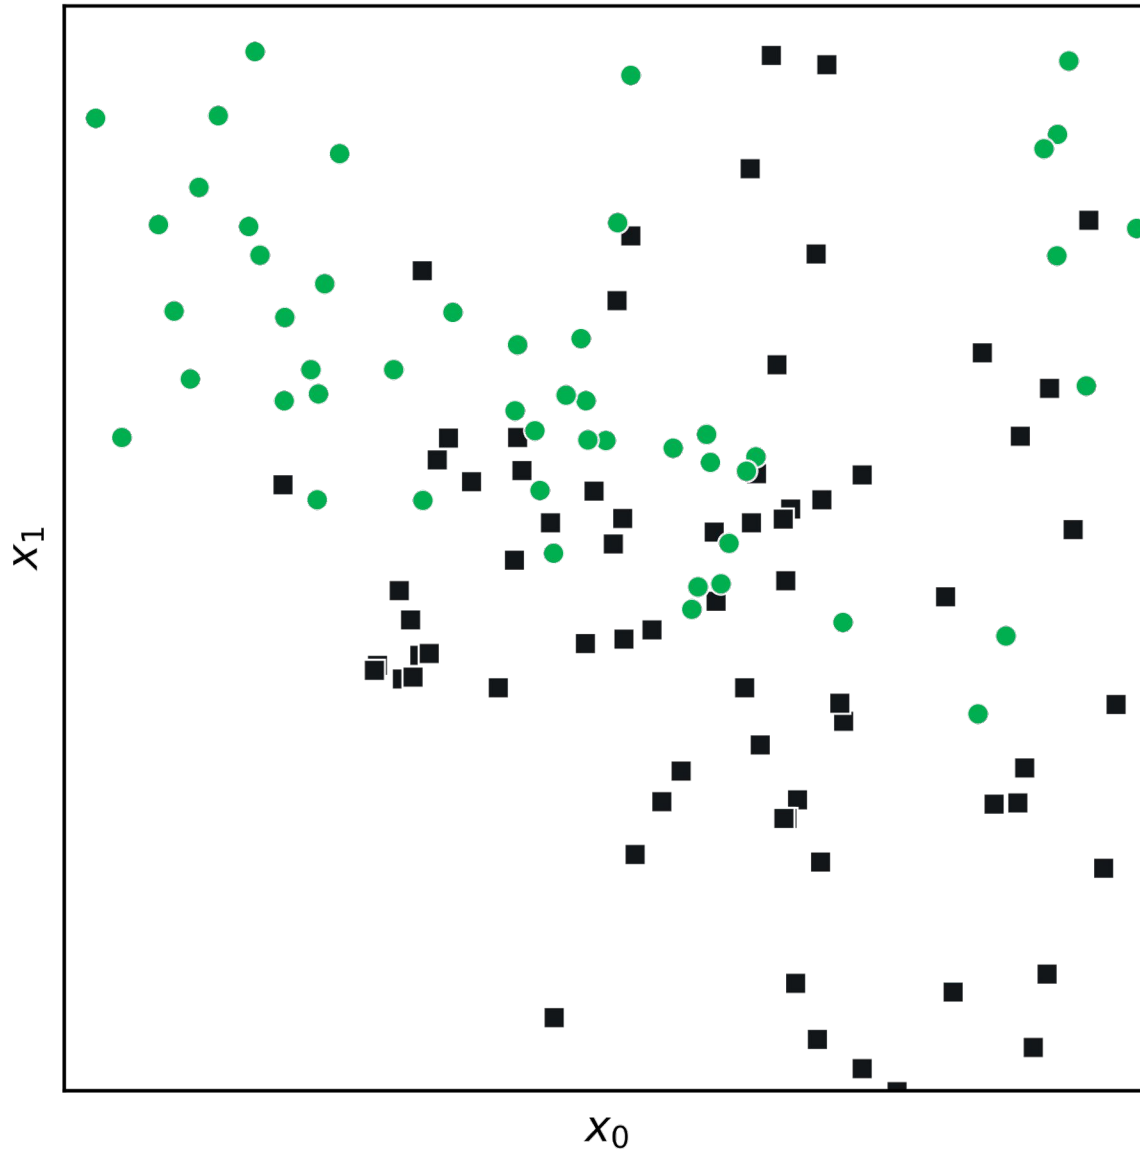
lower bound on  
generalization error

irreducible error inherent to the  
problem  
(e.g. you cannot predict the outcome of a flip of a  
fair coin any more than 50% of the time)

# Bias-Variance Tradeoff

$$\text{generalization error} = \text{bias}^2 + \text{variance} + \text{noise}$$

# Classification feature space



# What's the best we can do for binary classification?

If we know the probability distribution of the data

The Bayes decision rule

# Bayes' Rule

$$P(C|X) = \frac{P(X|C)P(C)}{P(X)}$$

Posterior

Likelihood Prior

Evidence

$X$  Features

$C$  Class label

i.e.  $C \in \{c_0, c_1\}$  for the binary case

## Bayes' Decision Rule:

choose the most probable class given the data

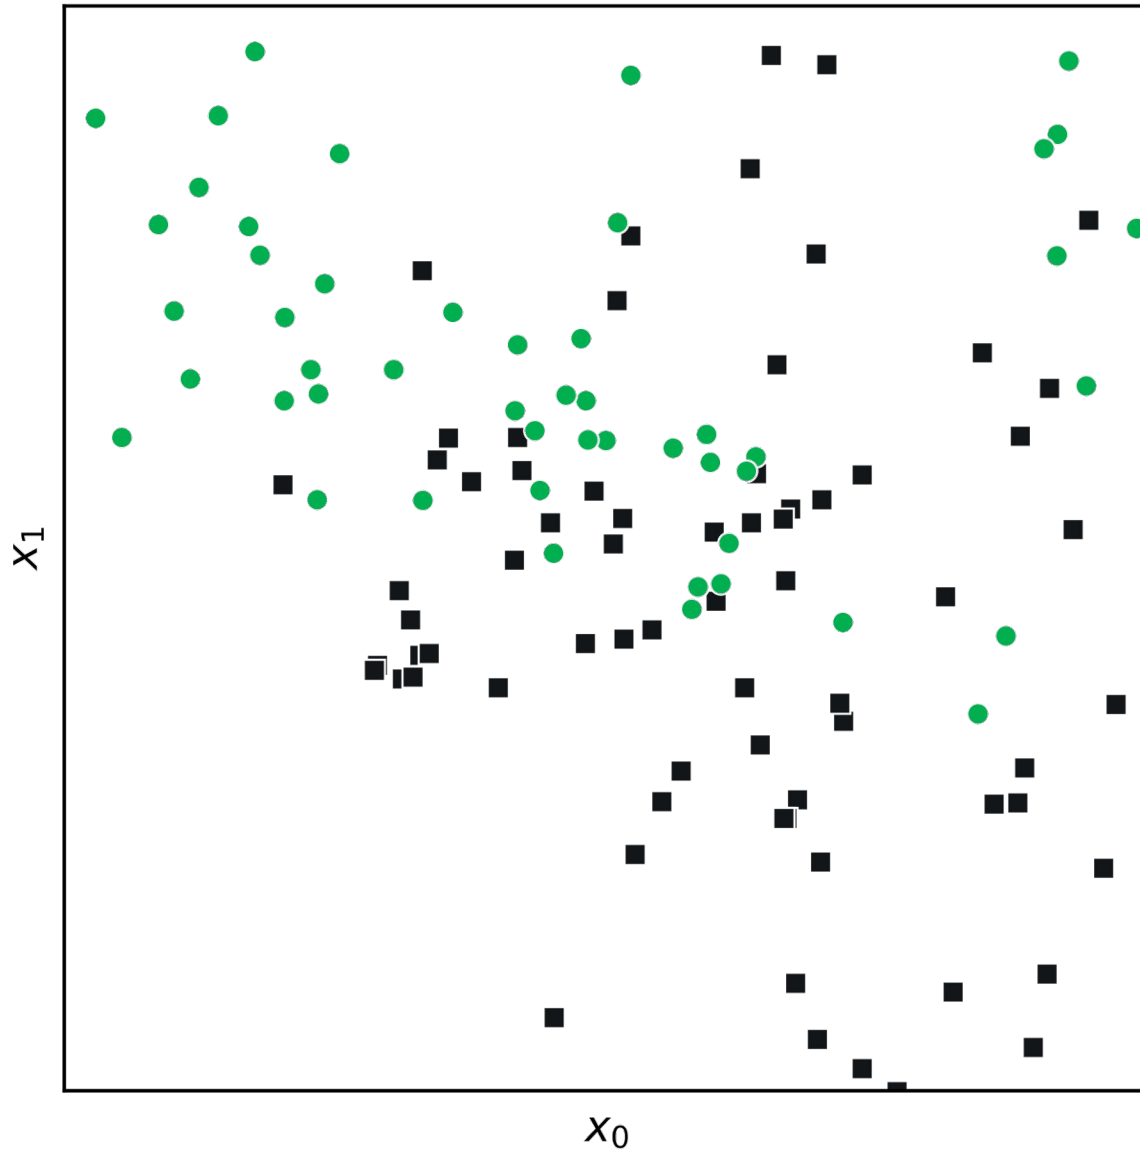
If  $P(C_i = c_1 | X_i) > P(C_i = c_0 | X_i)$  then  $\hat{y} = c_1$

otherwise

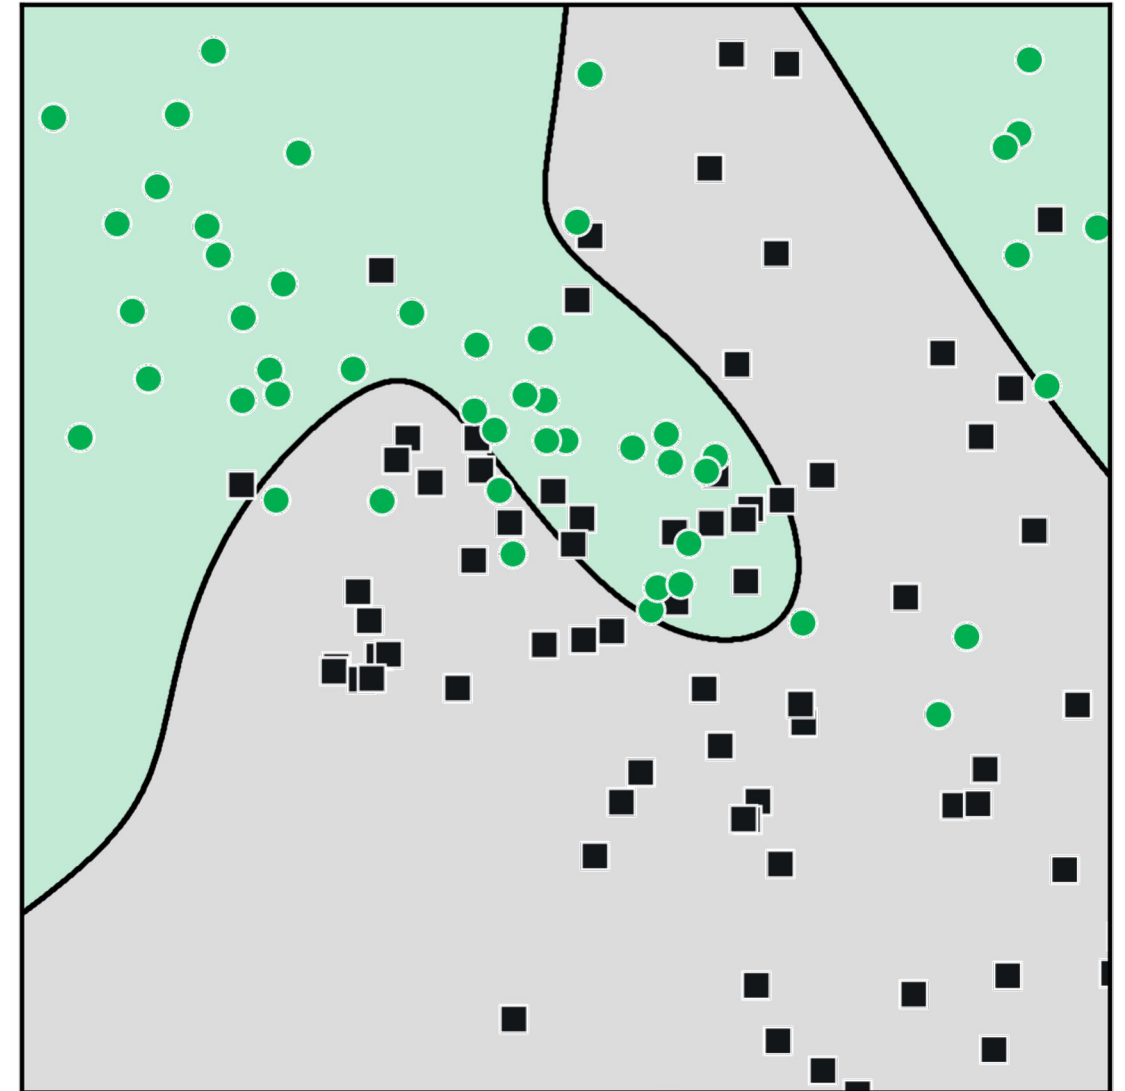
$\hat{y} = c_0$

- If the distributions are correct, this decision rule is **optimal**
- Rarely do we have enough information to use this in practice

# Classification feature space



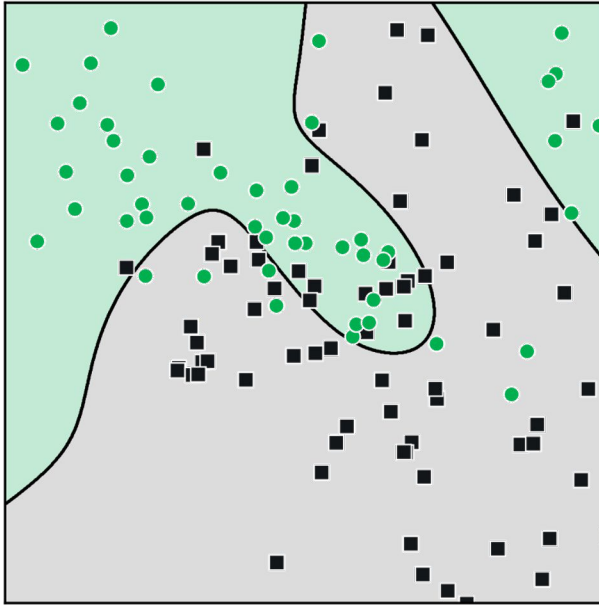
Bayes Decision Boundary



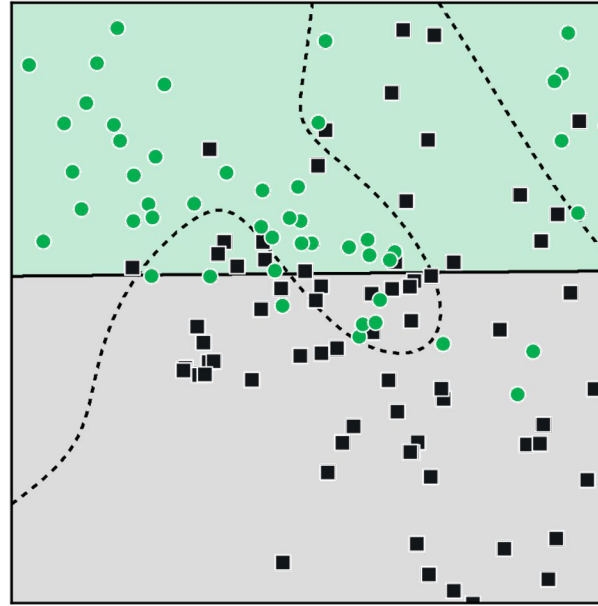


# Decision Boundary Examples

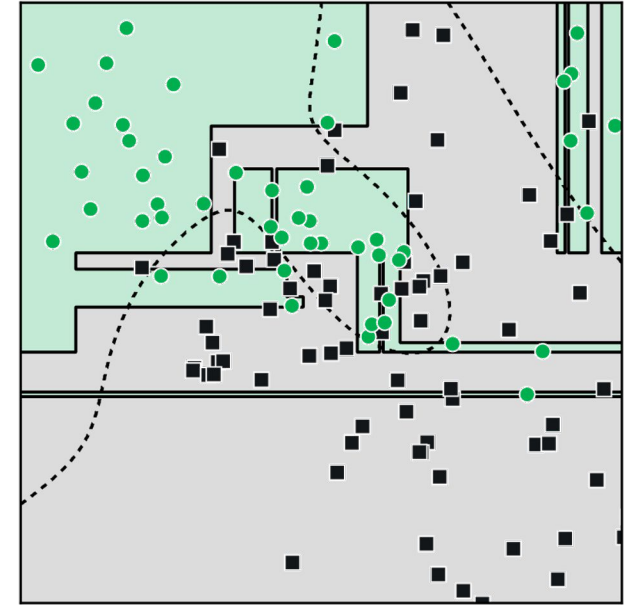
Bayes Decision Boundary



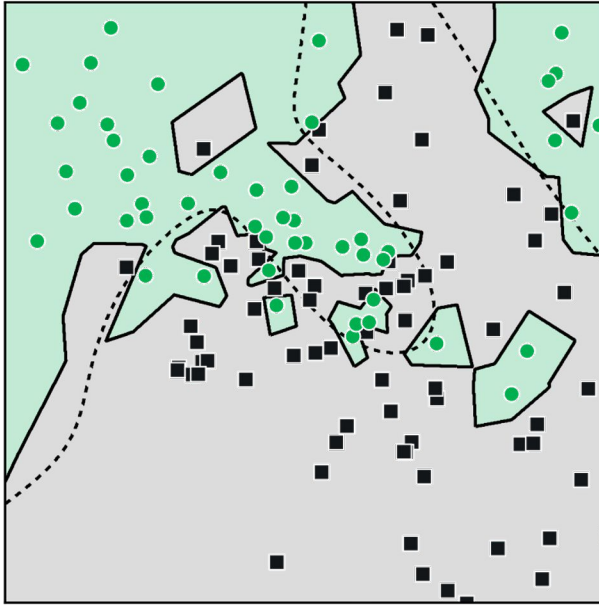
Linear Classifier



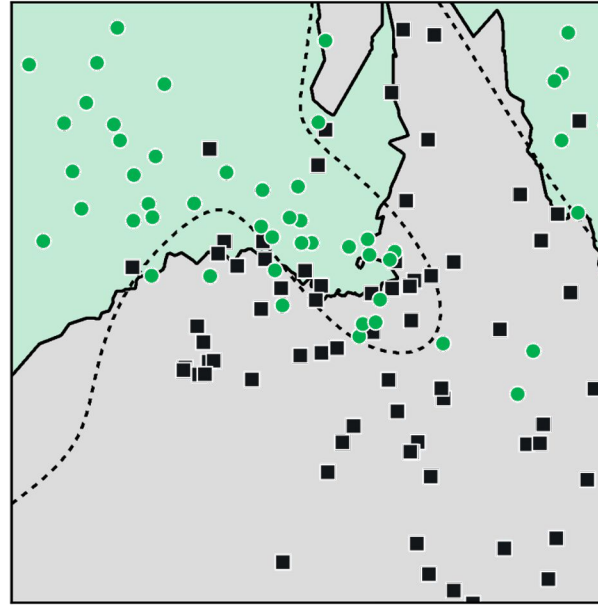
Decision Tree



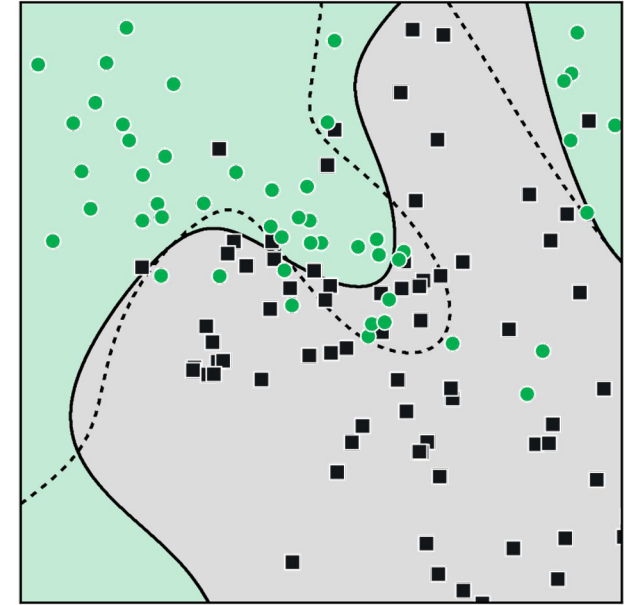
k=1 Nearest Neighbor



k=15 Nearest Neighbor

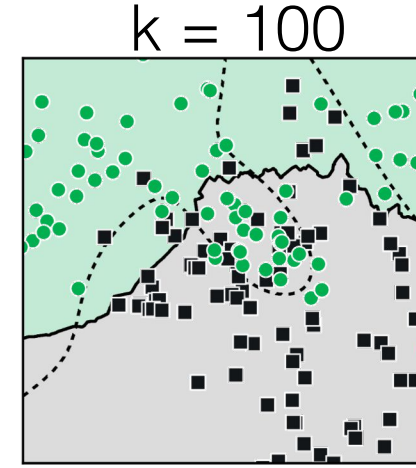
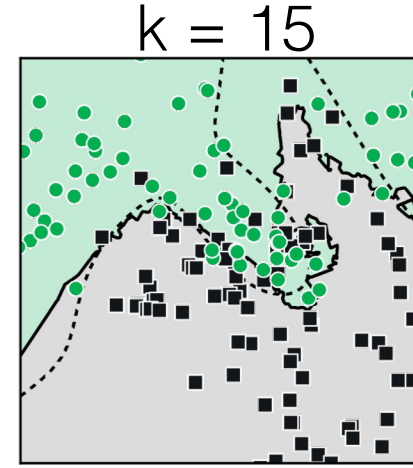
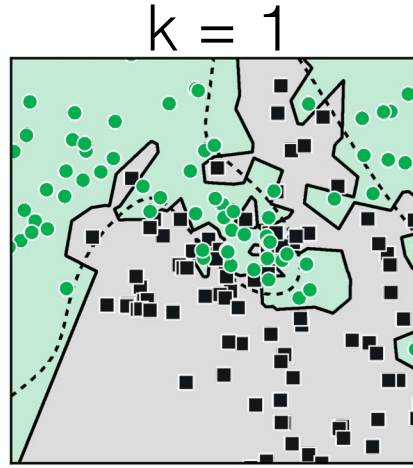


Support Vector Machine



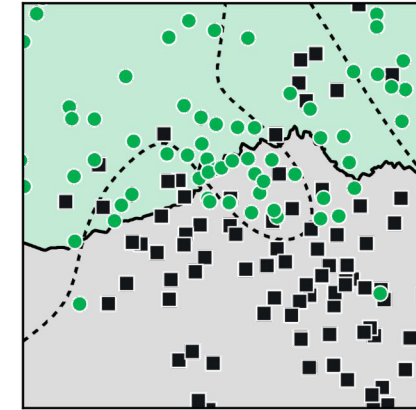
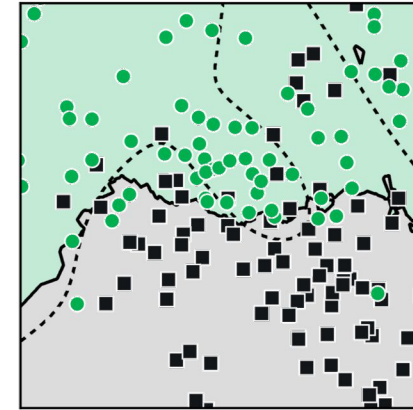
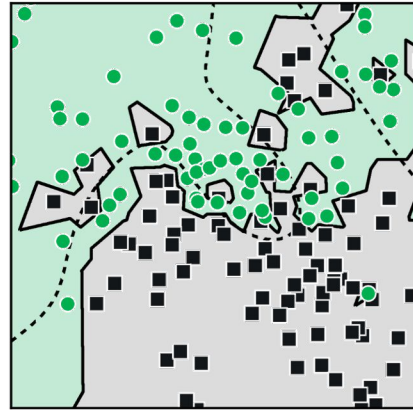
# Bias Variance Tradeoff

Sample 1



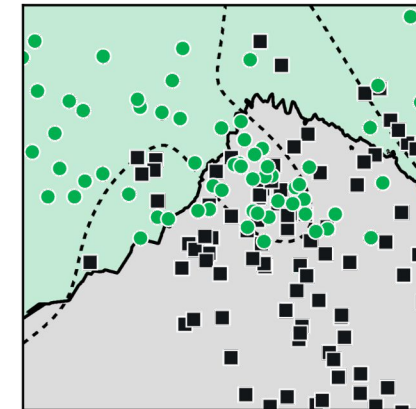
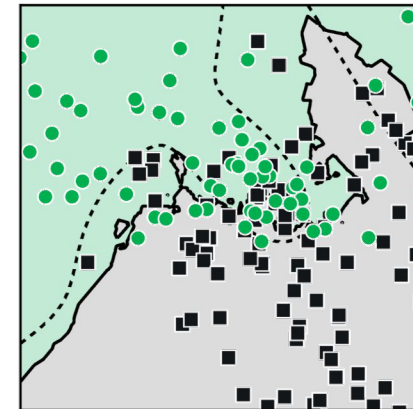
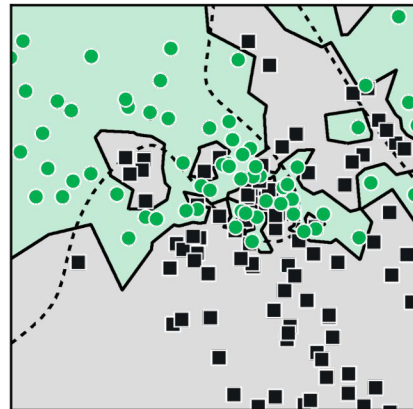
higher bias  
underfit

Sample 2

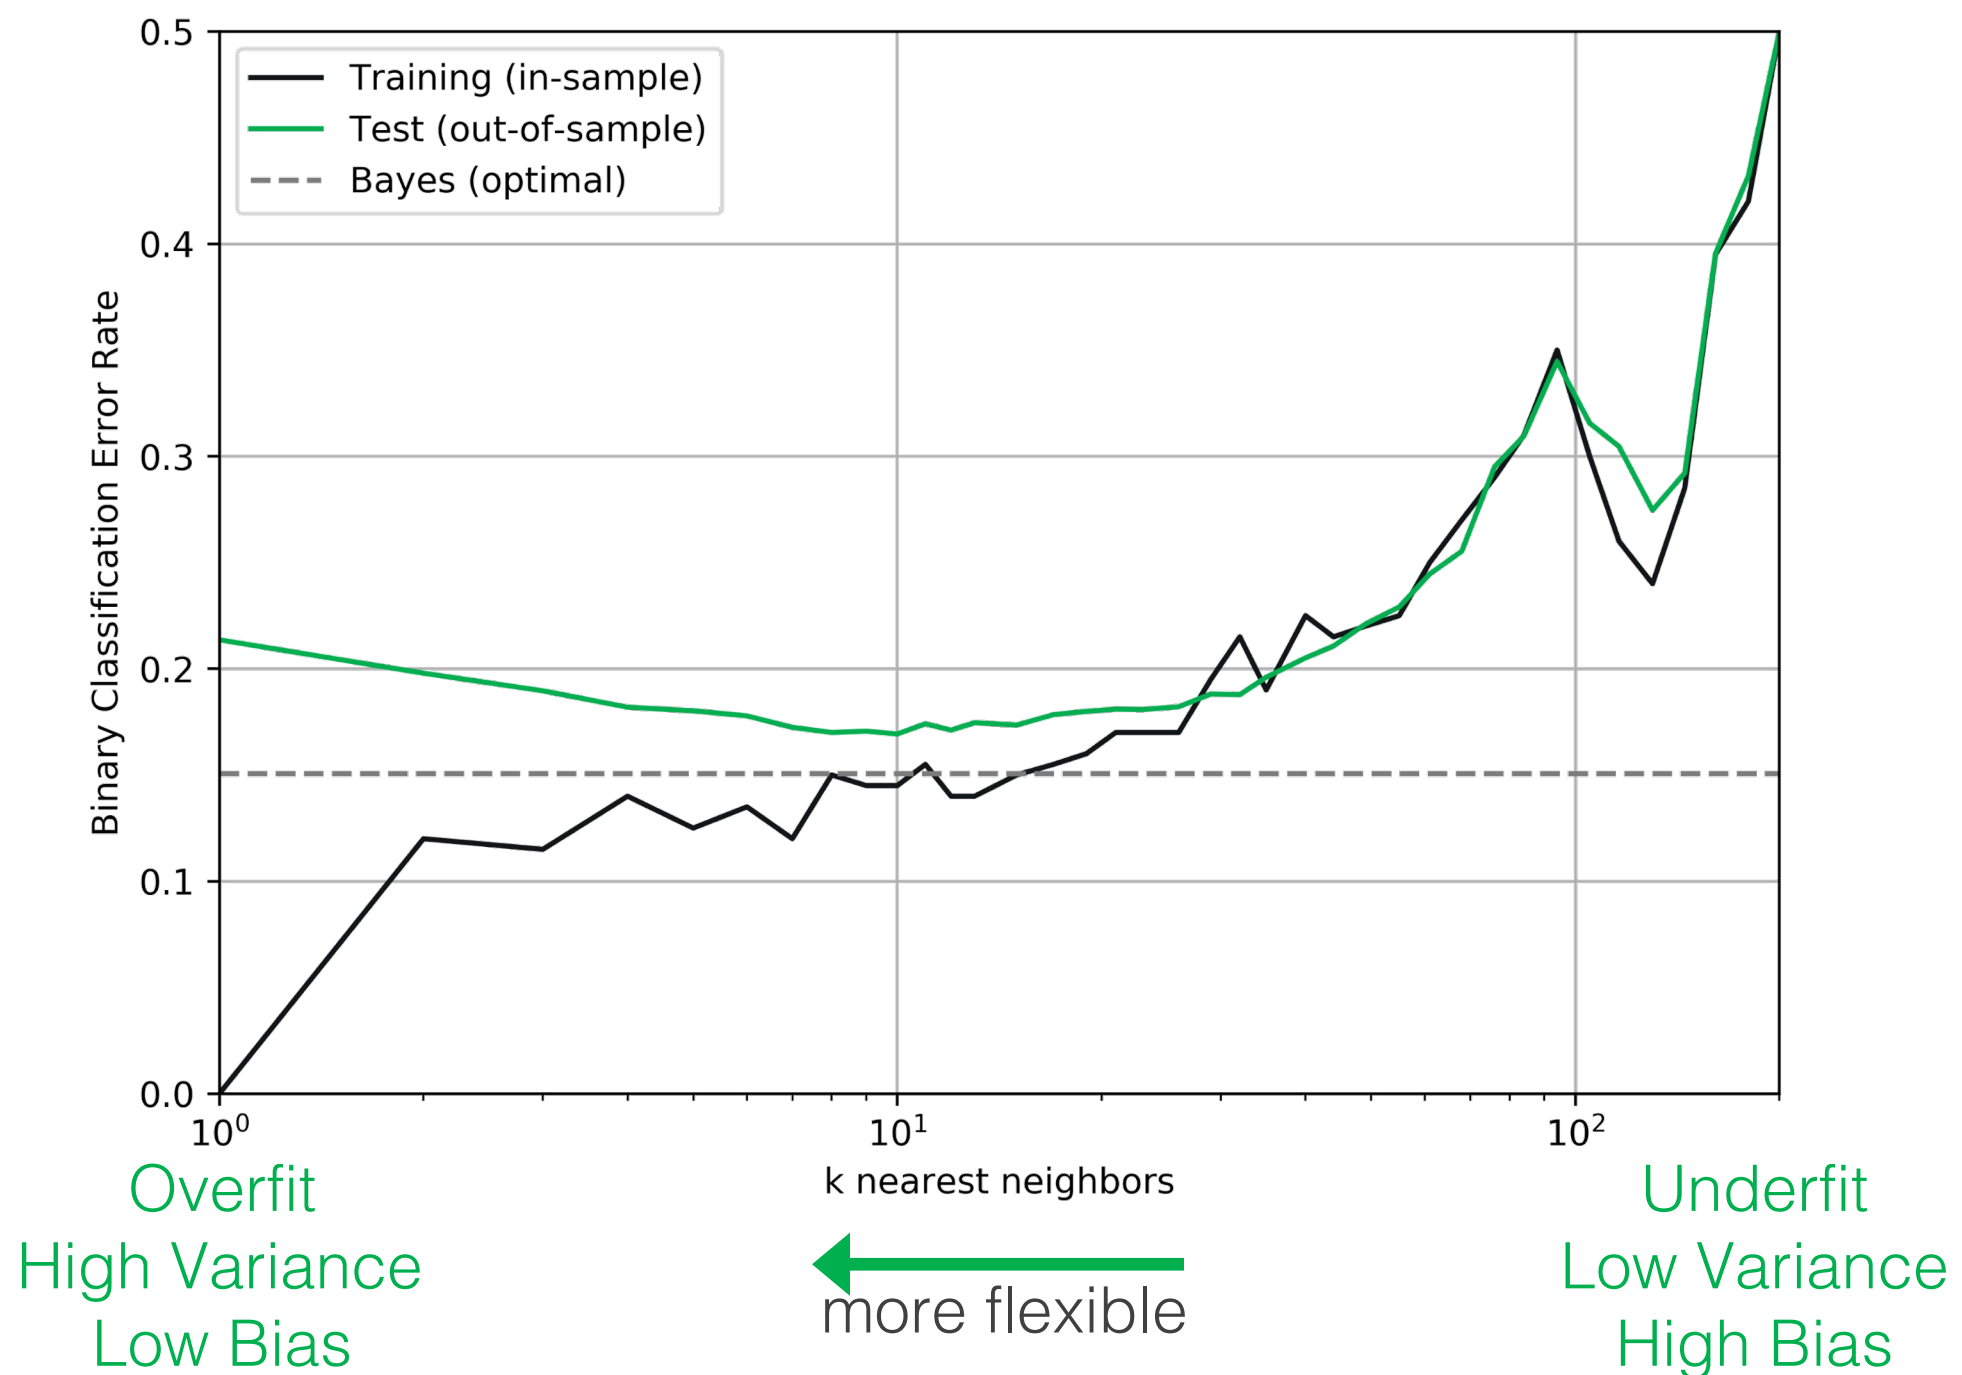


higher variance  
overfit

Sample 3



# Bias Variance Tradeoff

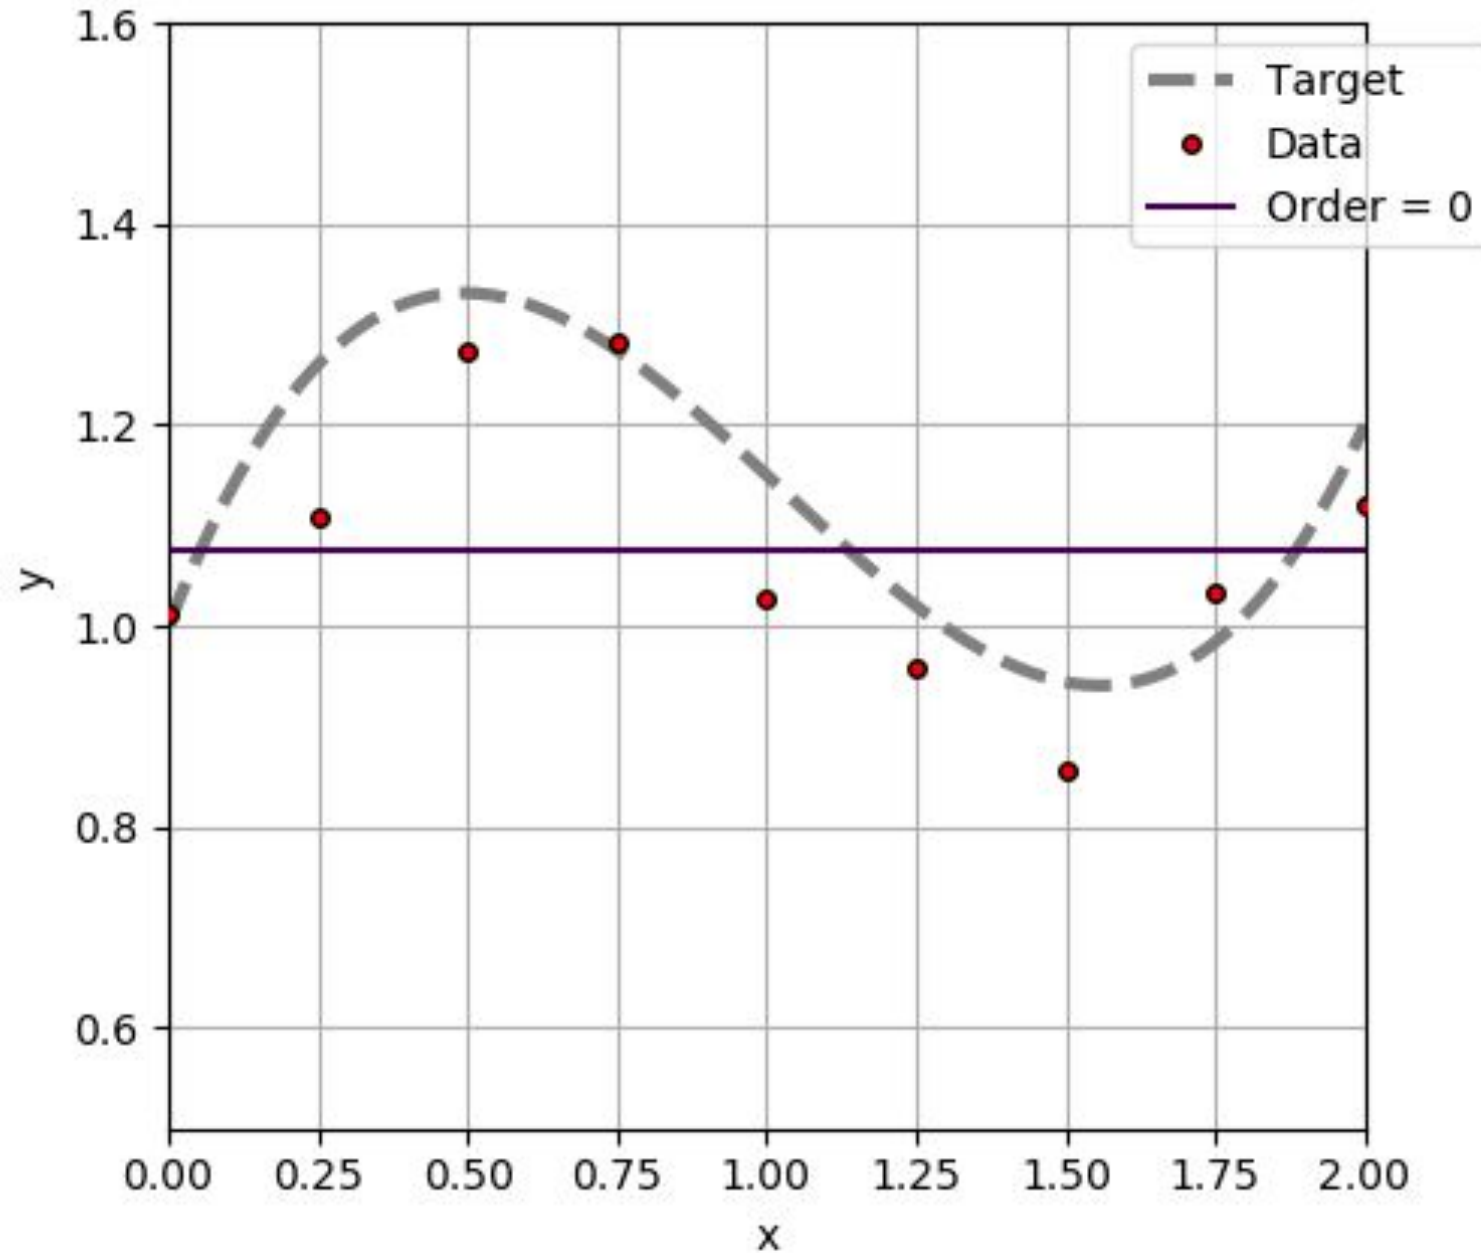


**This tradeoff is equally challenging for regression**

# Linear Regression

$$\hat{y}_i = \sum_{j=0}^m a_j x_i^j$$

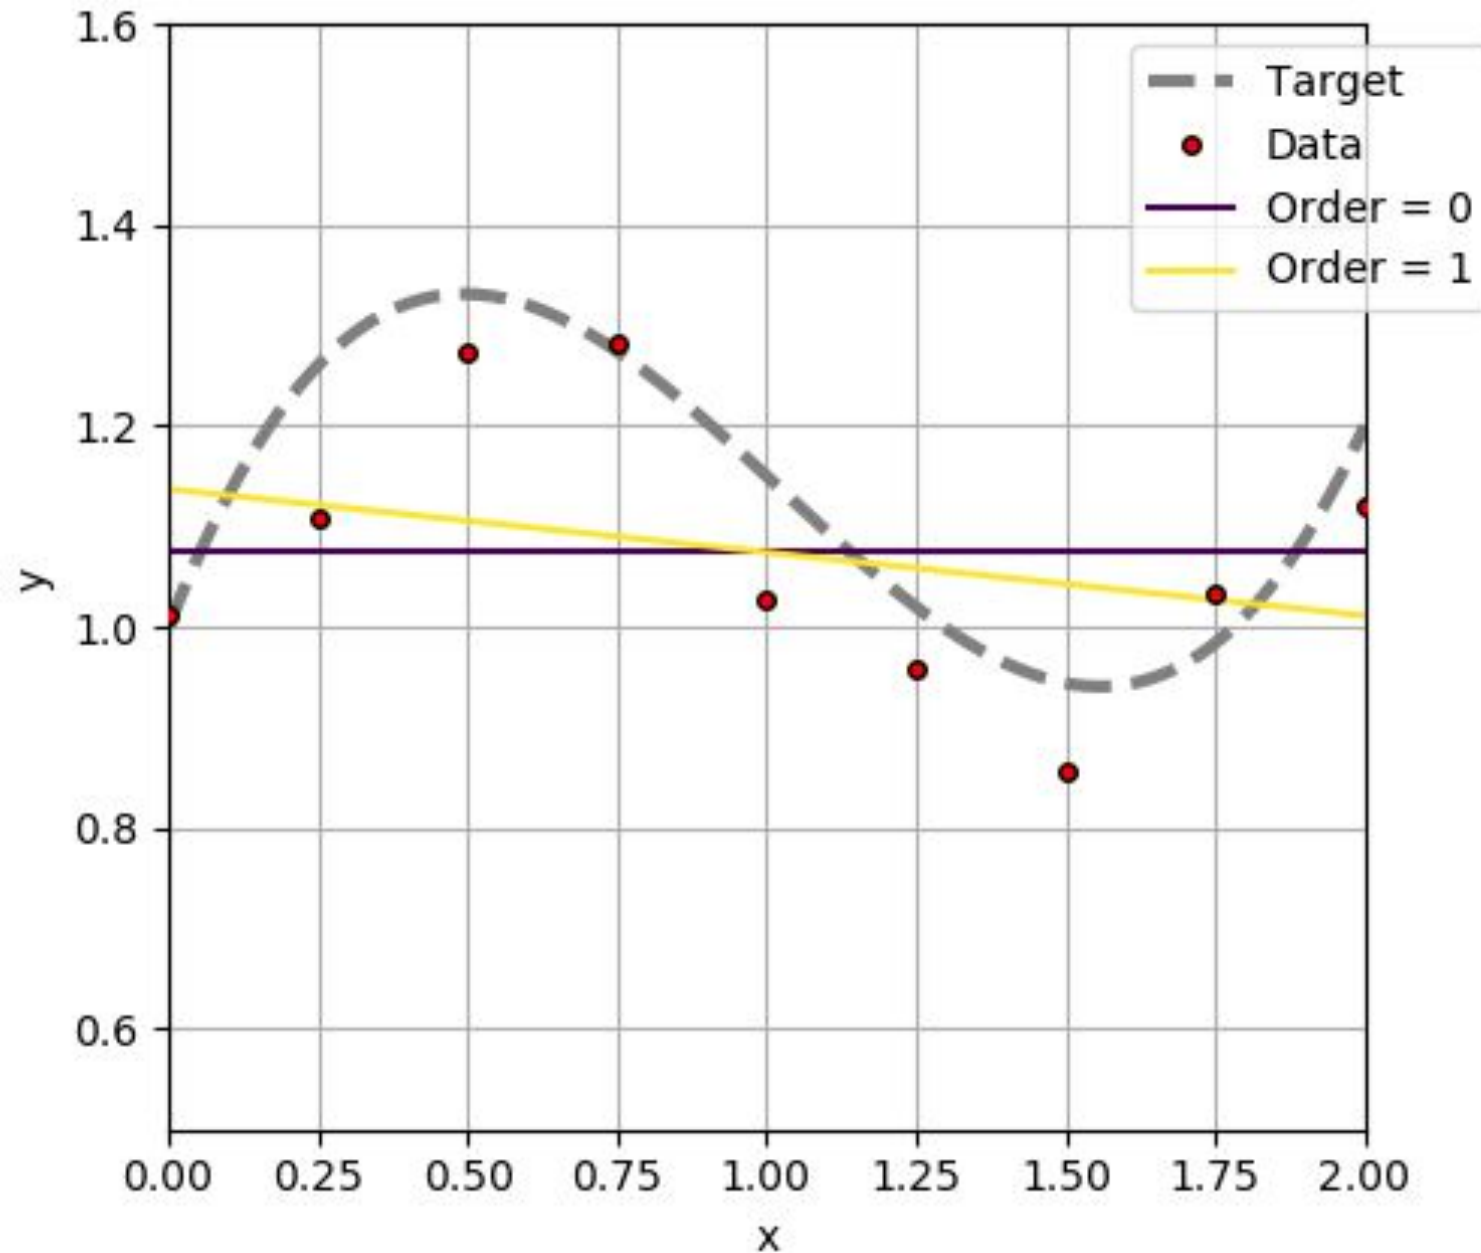
m is the model order



# Linear Regression

$$\hat{y}_i = \sum_{j=0}^m a_j x_i^j$$

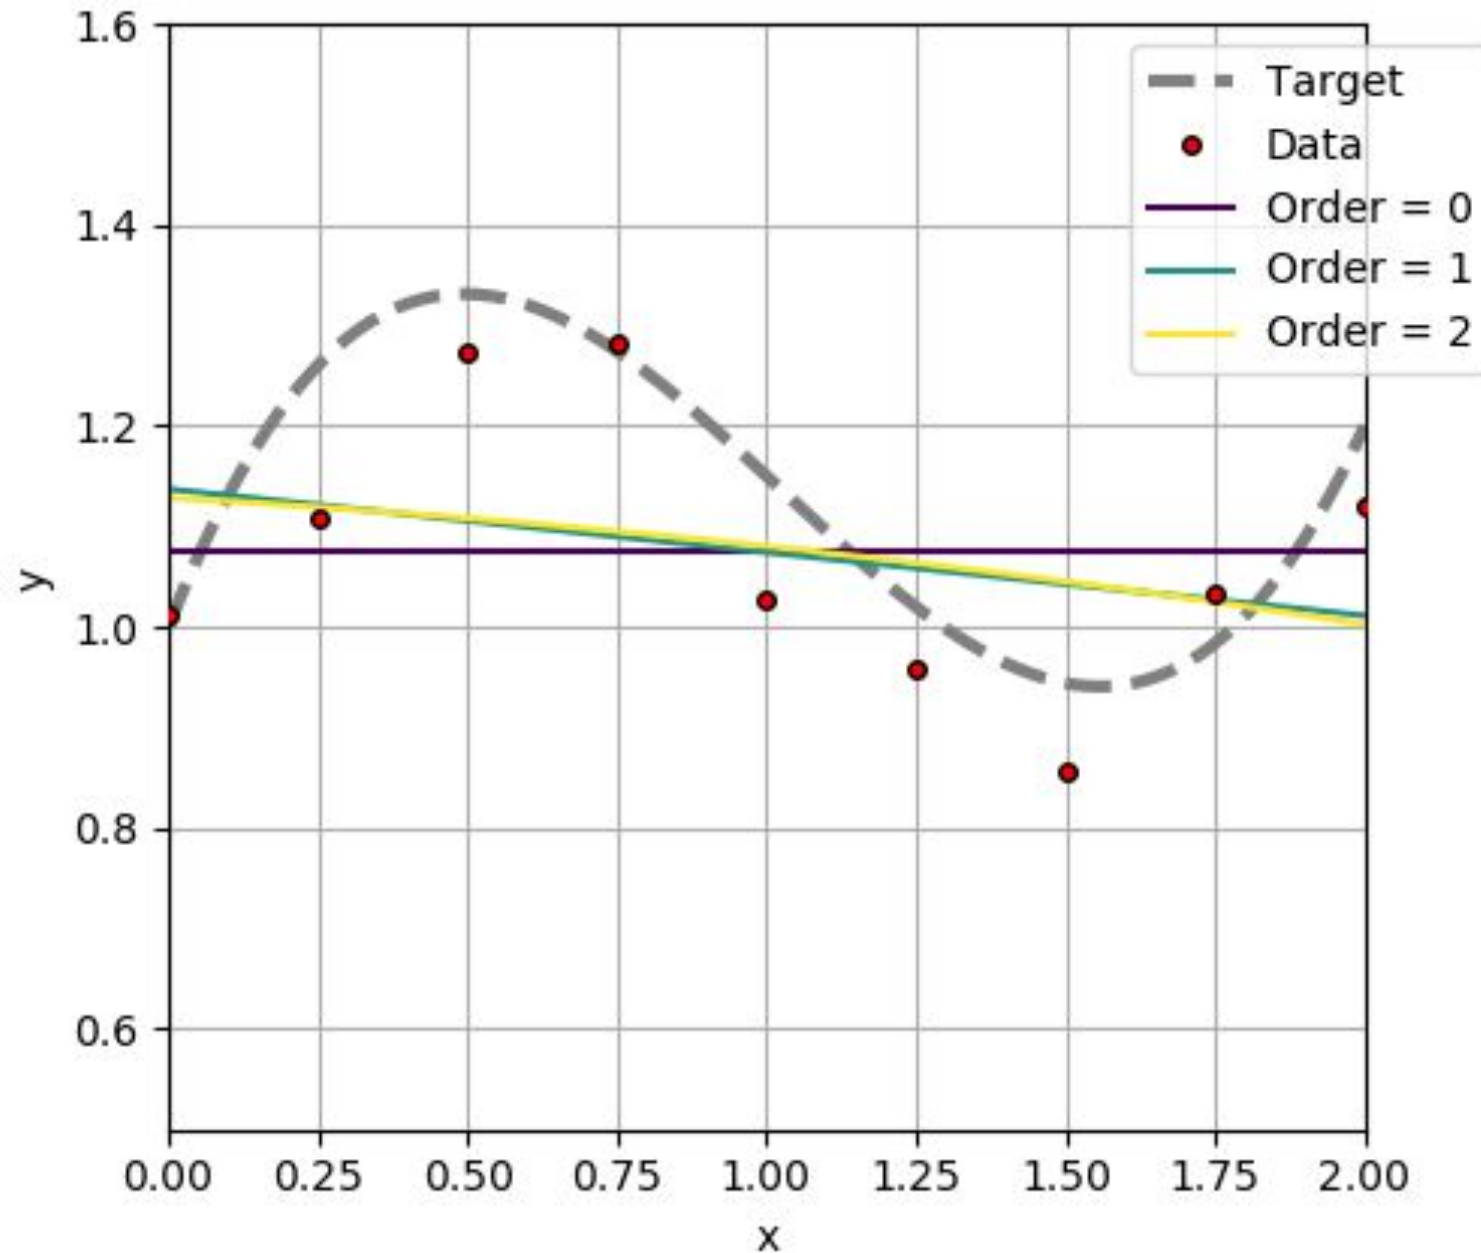
m is the model order



# Linear Regression

$$\hat{y}_i = \sum_{j=0}^m a_j x_i^j$$

m is the model order

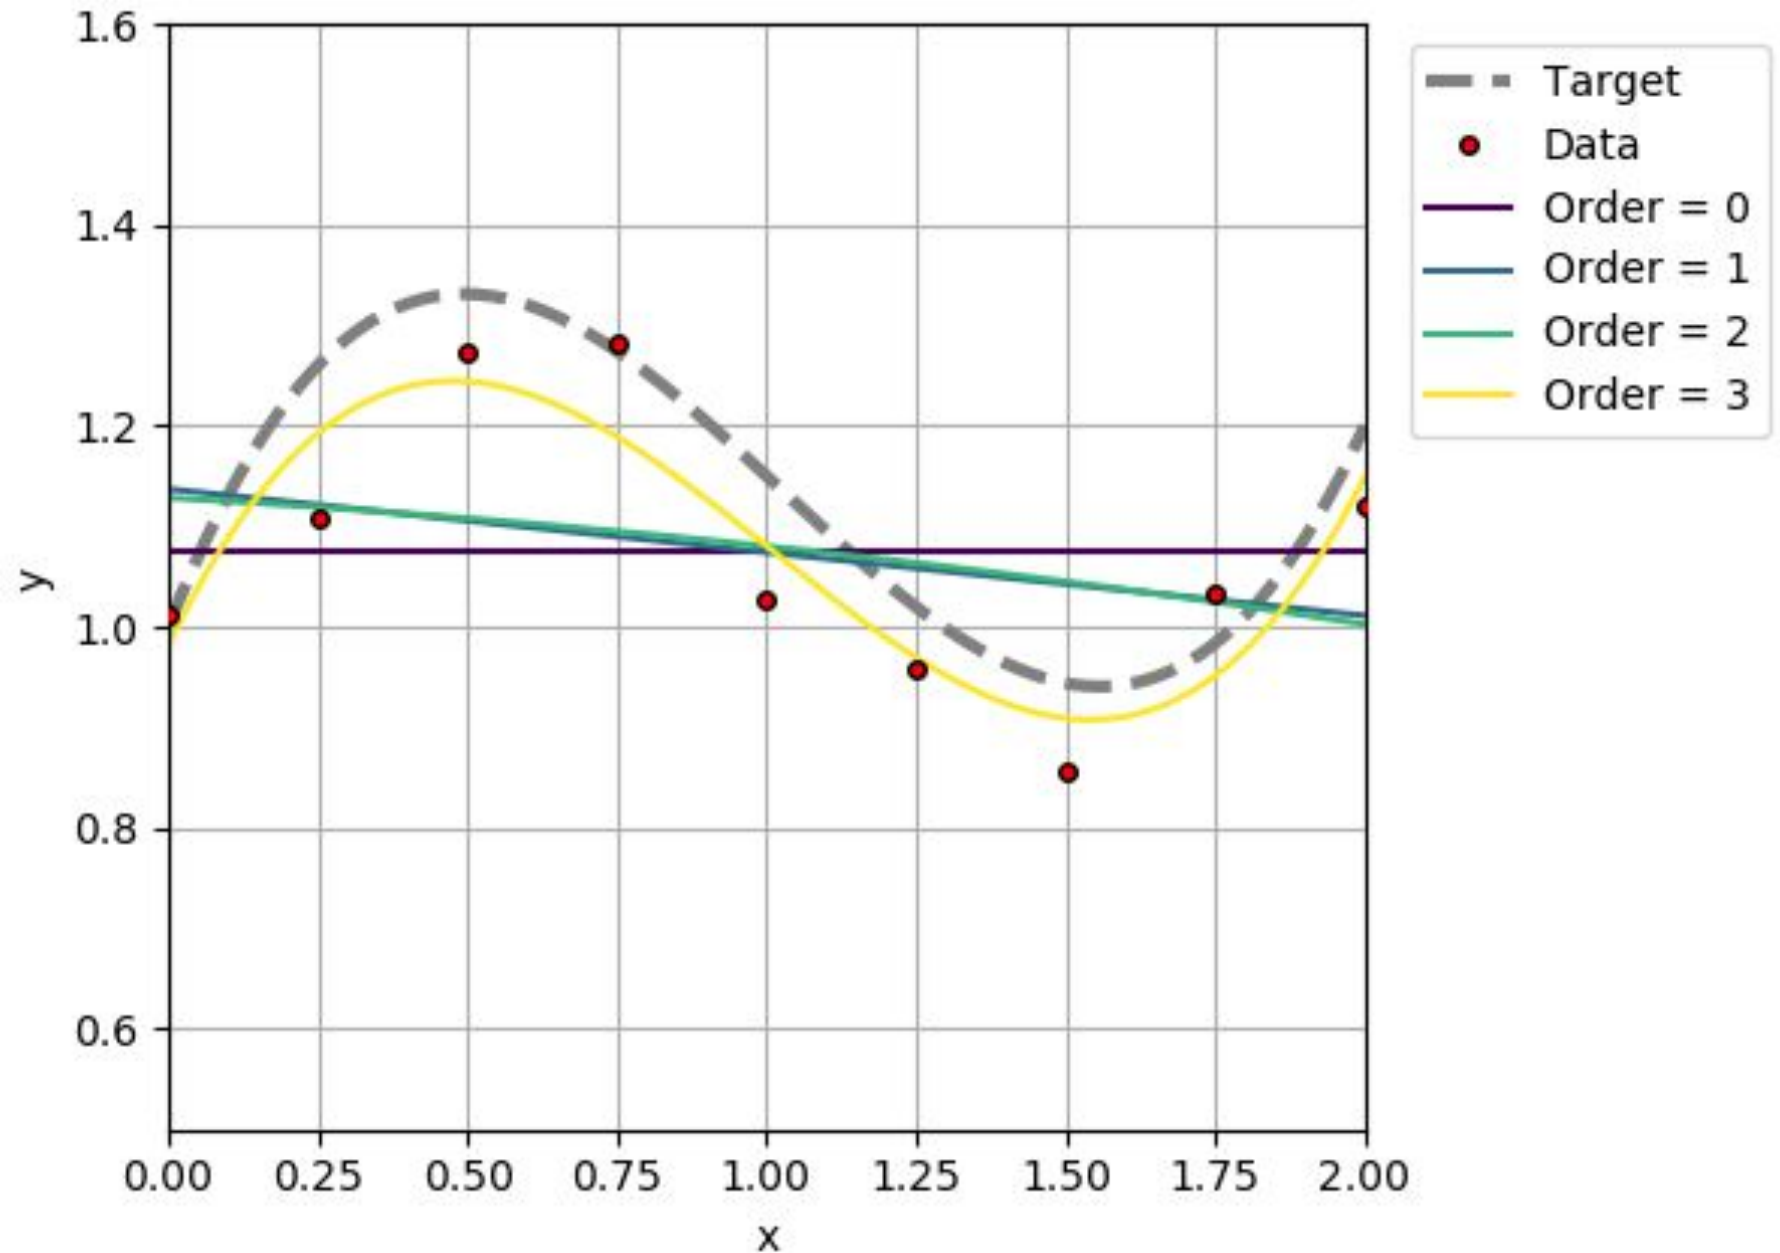




# Linear Regression

$$\hat{y}_i = \sum_{j=0}^m a_j x_i^j$$

m is the model order

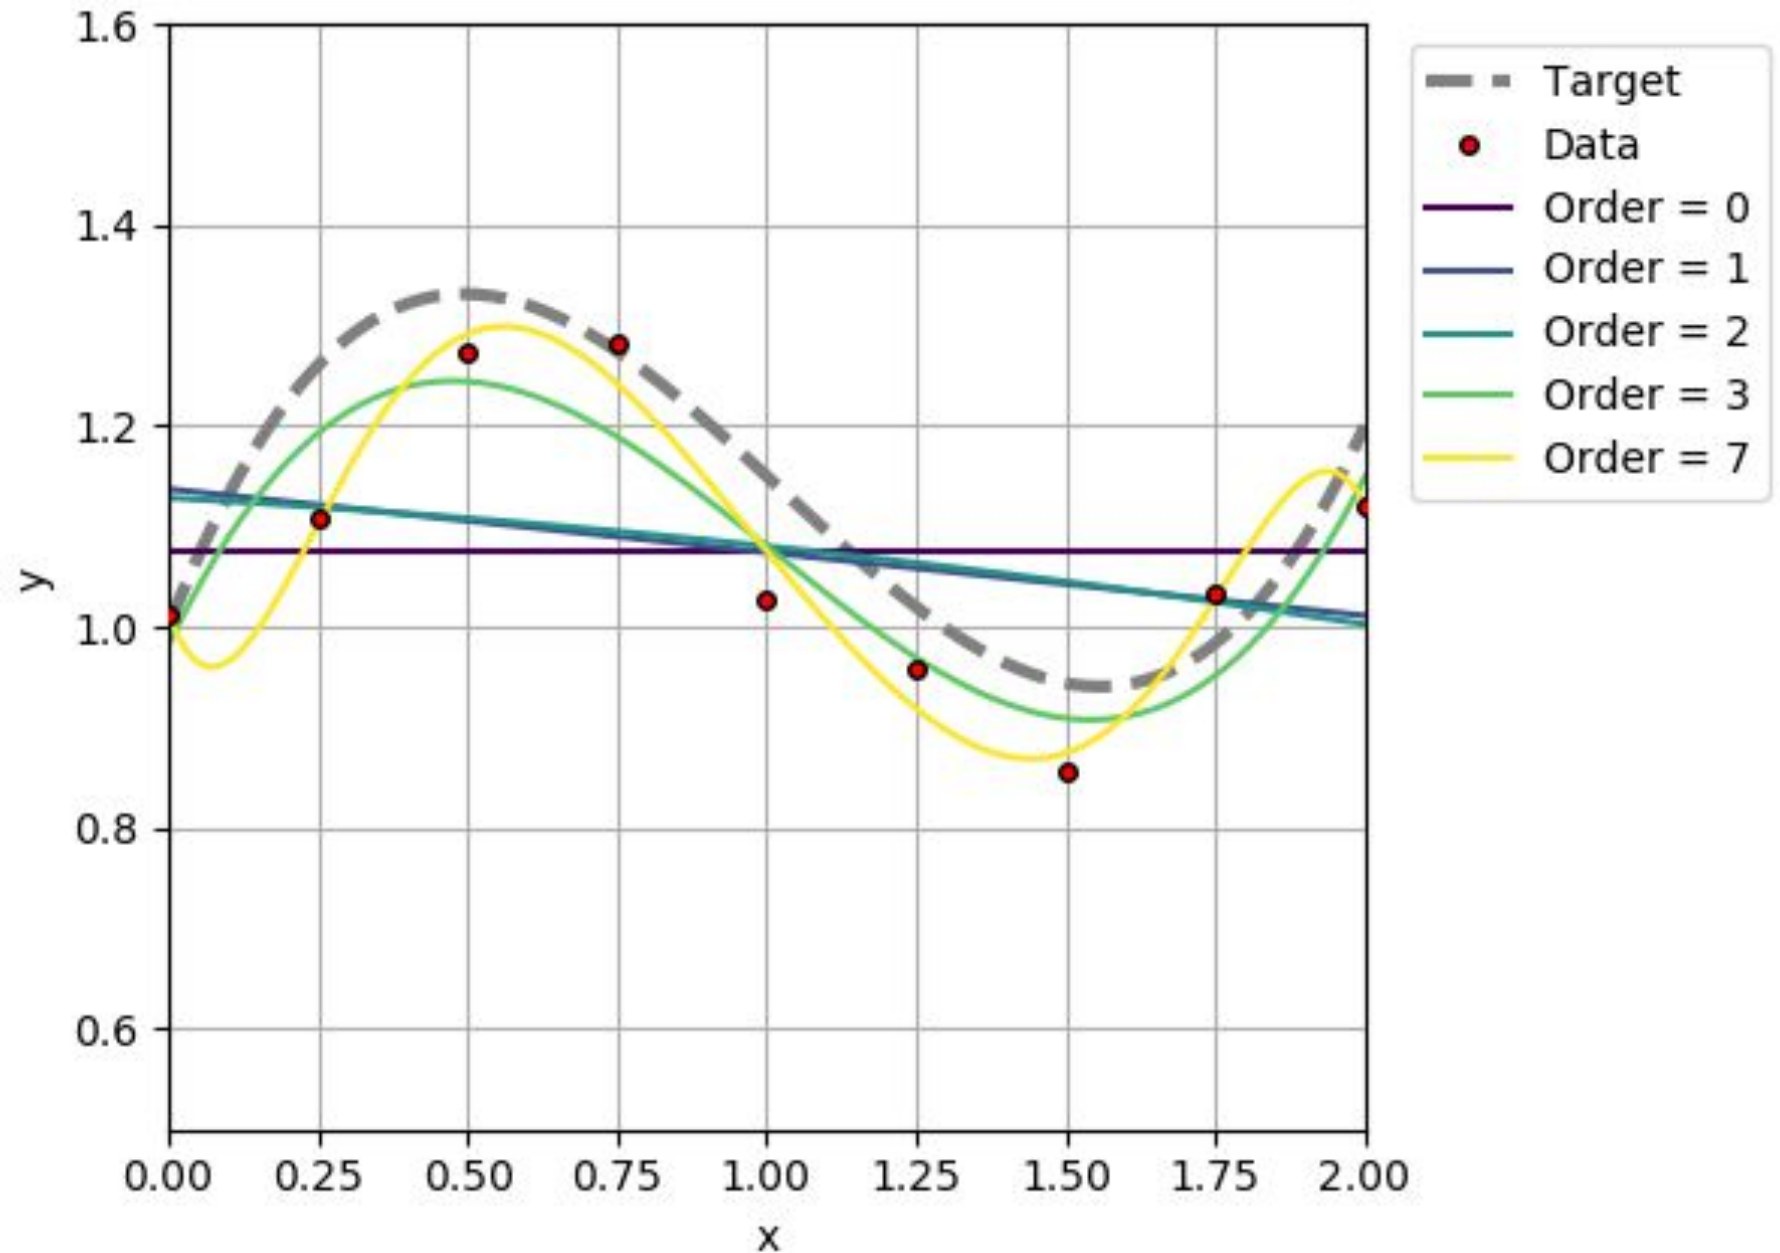




# Linear Regression

$$\hat{y}_i = \sum_{j=0}^m a_j x_i^j$$

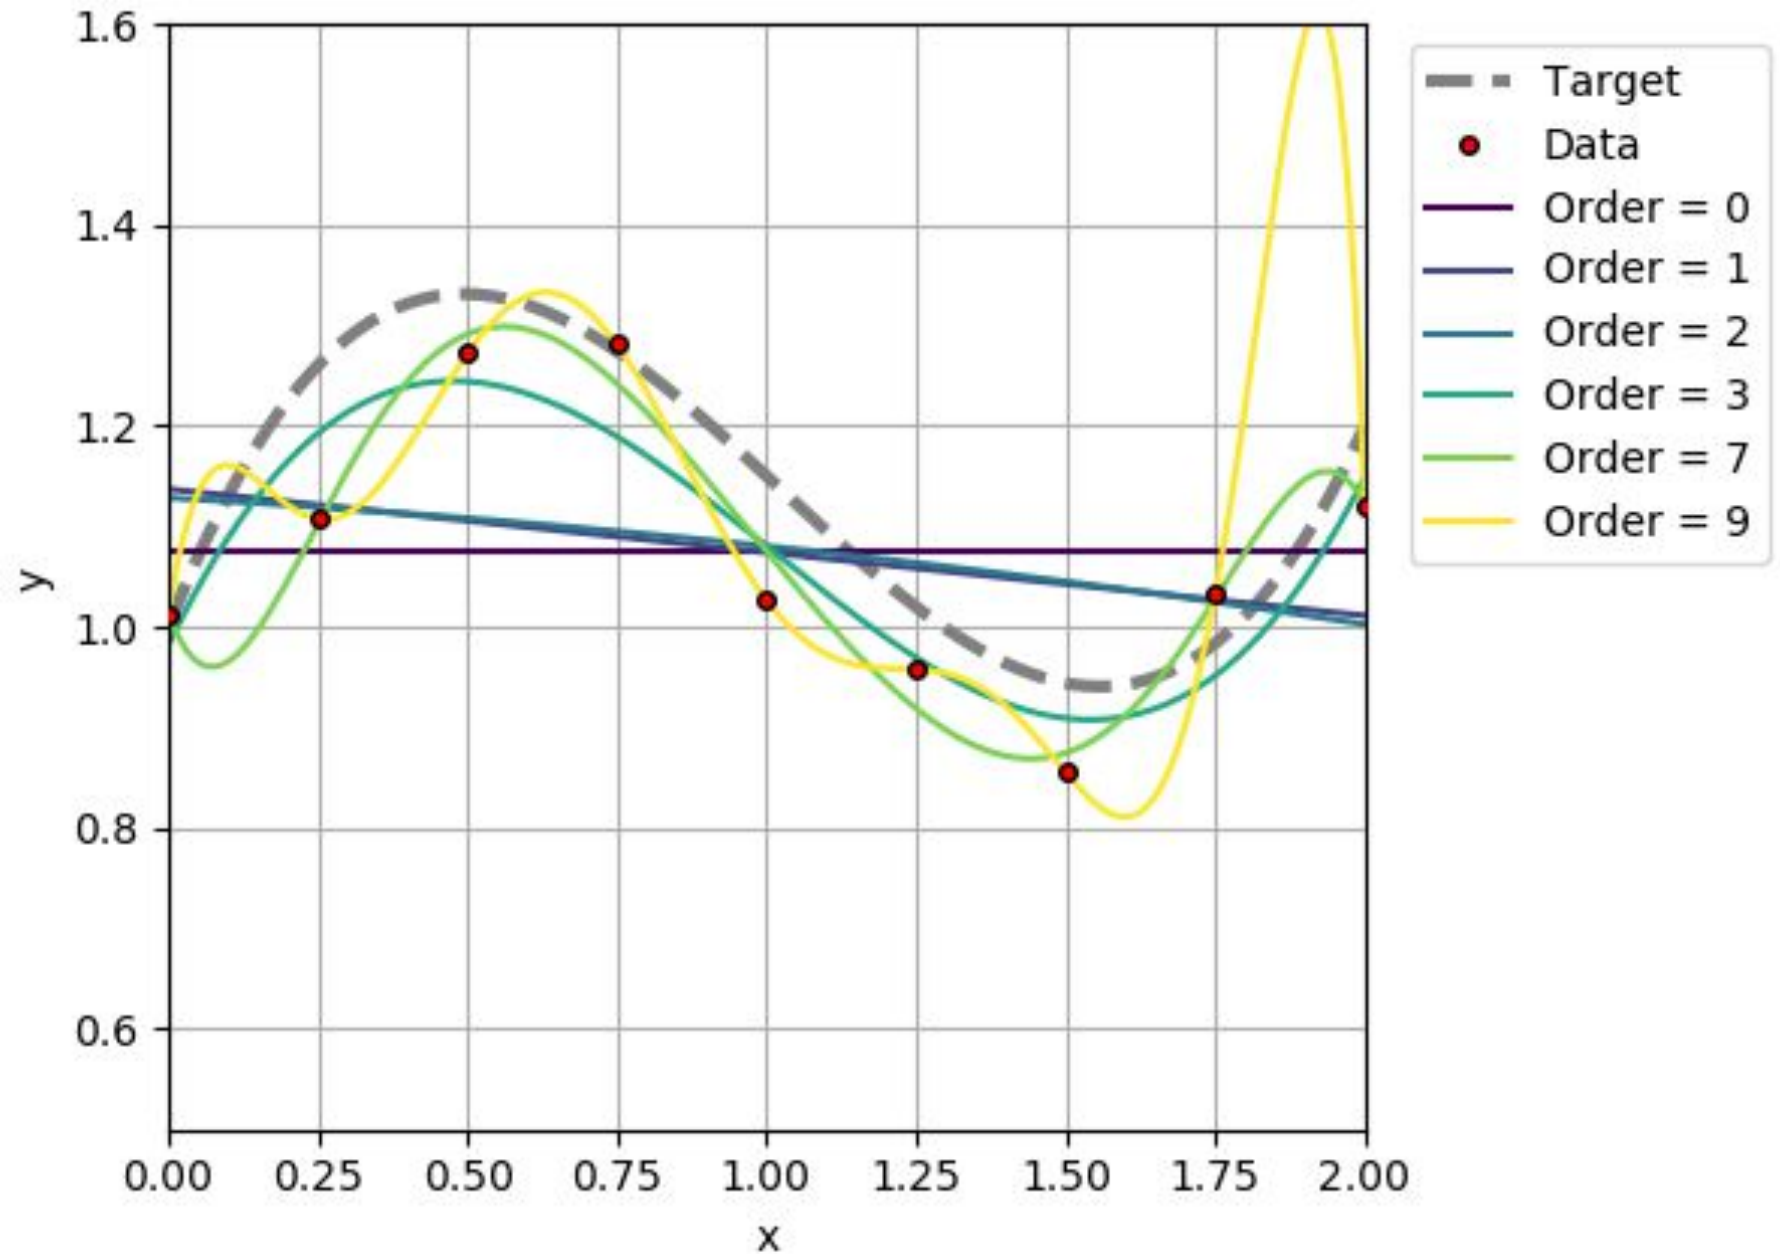
m is the model order



# Linear Regression

$$\hat{y}_i = \sum_{j=0}^m a_j x_i^j$$

m is the model order



# Problem

Too much flexibility leads to **overfit**

Too little flexibility leads to **underfit**

Over/underfit **hurts generalization** performance

## Solutions for overfitting

1. Add **more data** for training
2. Constrain model flexibility through **regularization**
3. Use model **ensembles**